



# Topology aware Cartesian grid mapping with MPI

## Issue 120 Reading

Christoph Niethammer    Rolf Rabenseifner

niethammer@hls.de

rabenseifner@hls.de

High Performance Computing Center (HLRS), University of Stuttgart, Germany

**HLRS, Stuttgart, March 02, 2019**

MPI Forum March 4-7, 2019, Chattanooga

For further information, see also <https://fs.hls.de/projects/par/mpi/EuroMPI2018-Cartesian/>



2018  
Höchstleistungsrechenzentrum Stuttgart



# The Problems of MPI\_Dims\_create + MPI\_Cart\_create

- The factorization of a given amount of MPI processes must be
  - Application topology aware [1]
  - Hardware topology aware
- Current definition of MPI\_Dims\_create is not prepared for this
- Extreme differences in latency and accumulated bandwidth between **inter**-node and **intra**-node communication
- The reordering by MPI\_Cart\_create:
  - Many implementations do nothing [2]
  - A perfect reordering may require complex domain decomposition algorithms (e.g. Metis) [3]

} Slides 2-3

} Slides 4-6

} Slide 7-10

We propose a new interface together with a fast algorithm, which is application and hardware topology aware

} Slides 11-22

Implementation remarks

} Slide 23-24

Benchmark + results

} Slide 25-34

[1, 2, 3] see References on last slide

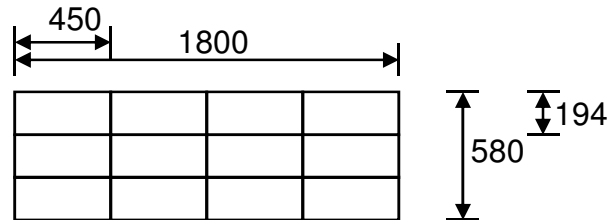
} Slide 35

# Examples

- Application topology awareness

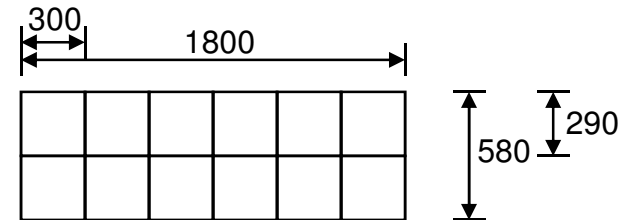
- 2-D example with 12 MPI processes and gridsize 1800x580

- **MPI\_Dims\_create** → 4x3



Boundary of a subdomain =  $2(450+194) = 1288$  😞

- **grid aware** → 6x2 processes

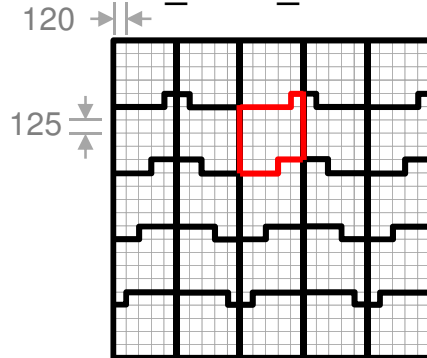


Boundary of a subdomain =  $2(300+290) = 1180$  😊

- Hardware topology awareness

- 2-D example with 25 nodes x 24 cores and gridsize 3000x3000

- **MPI\_Dims\_create** → 25 x 24

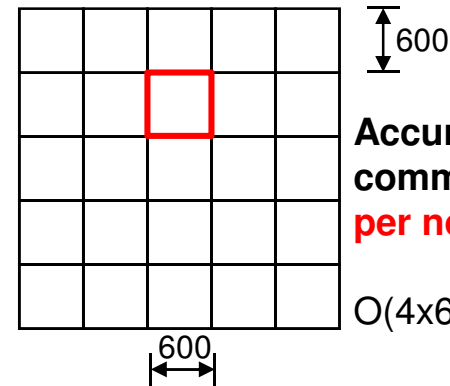


**Accumulated communication per node**

$O(10 \times 120 + 12 \times 125)$   
 $= O(2700)$  😞

- **Hardware aware**

→ (5 nodes x 6 cores) X (5 nodes x 4 cores)



**Accumulated communication per node**

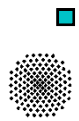
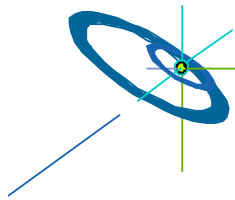
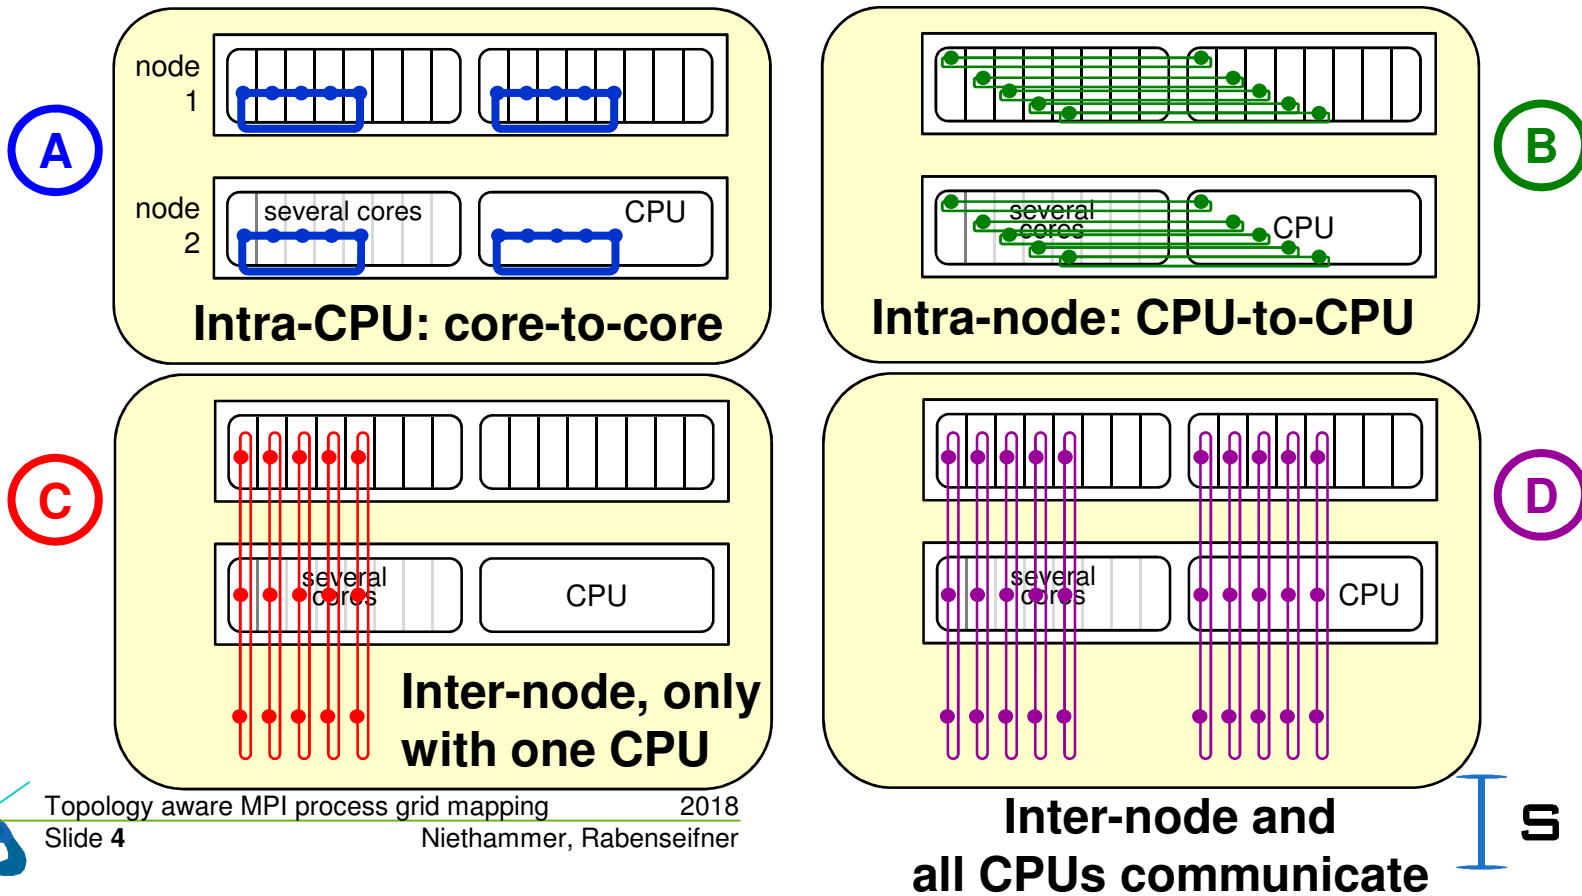
$O(4 \times 600) = O(2400)$  😊

# Ring Benchmarks for Inter- and Intra-node Communication

Benchmark `halo_irecv_send_multiplelinks_toggle.c`

- Varying message size,
- number of **communication cores per CPU**, and
- four communication schemes (example with 5 **communicating cores per CPU**)

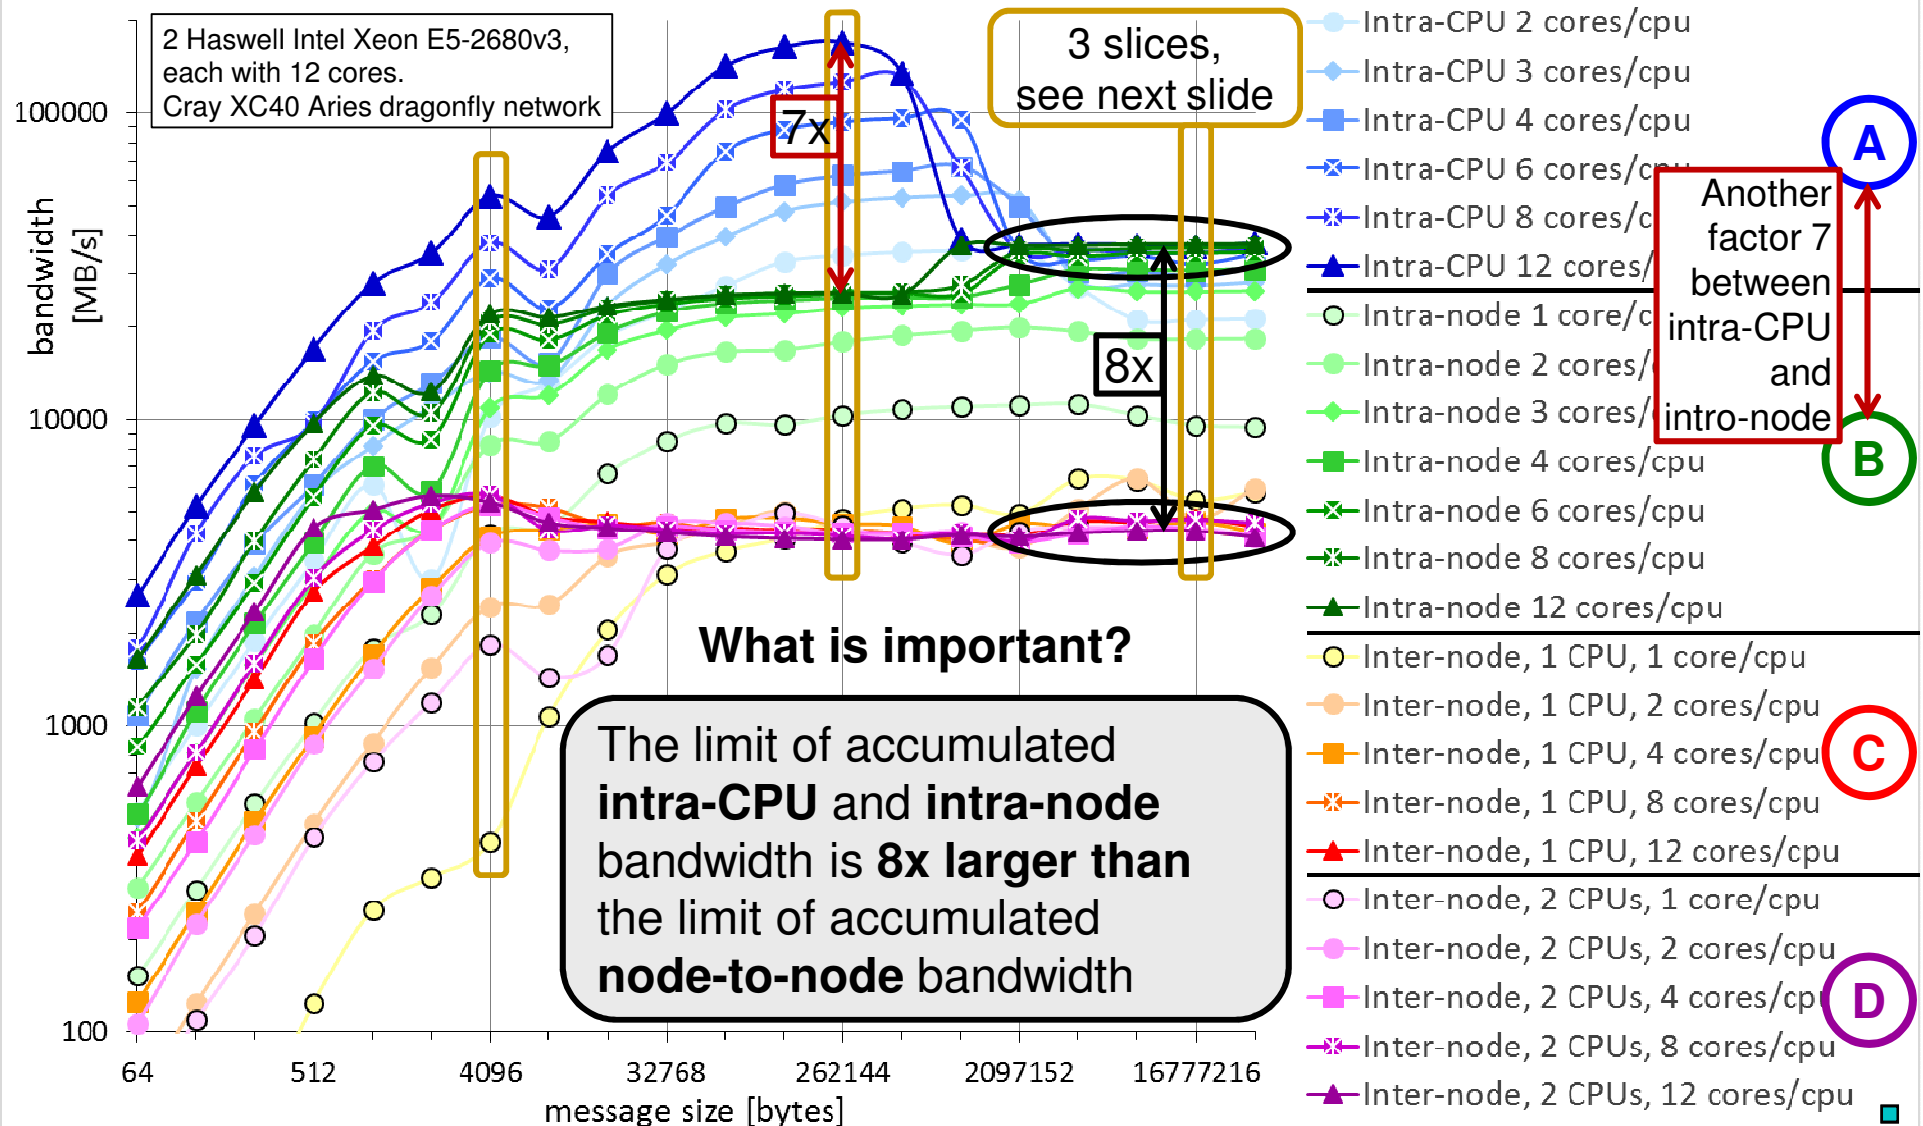
See HLRS online courses  
<http://www.hlrs.de/training/par-prog-ws/>  
→ Practical → MPI.tar.gz  
→ subdirectory MPI/course/C/1sided/



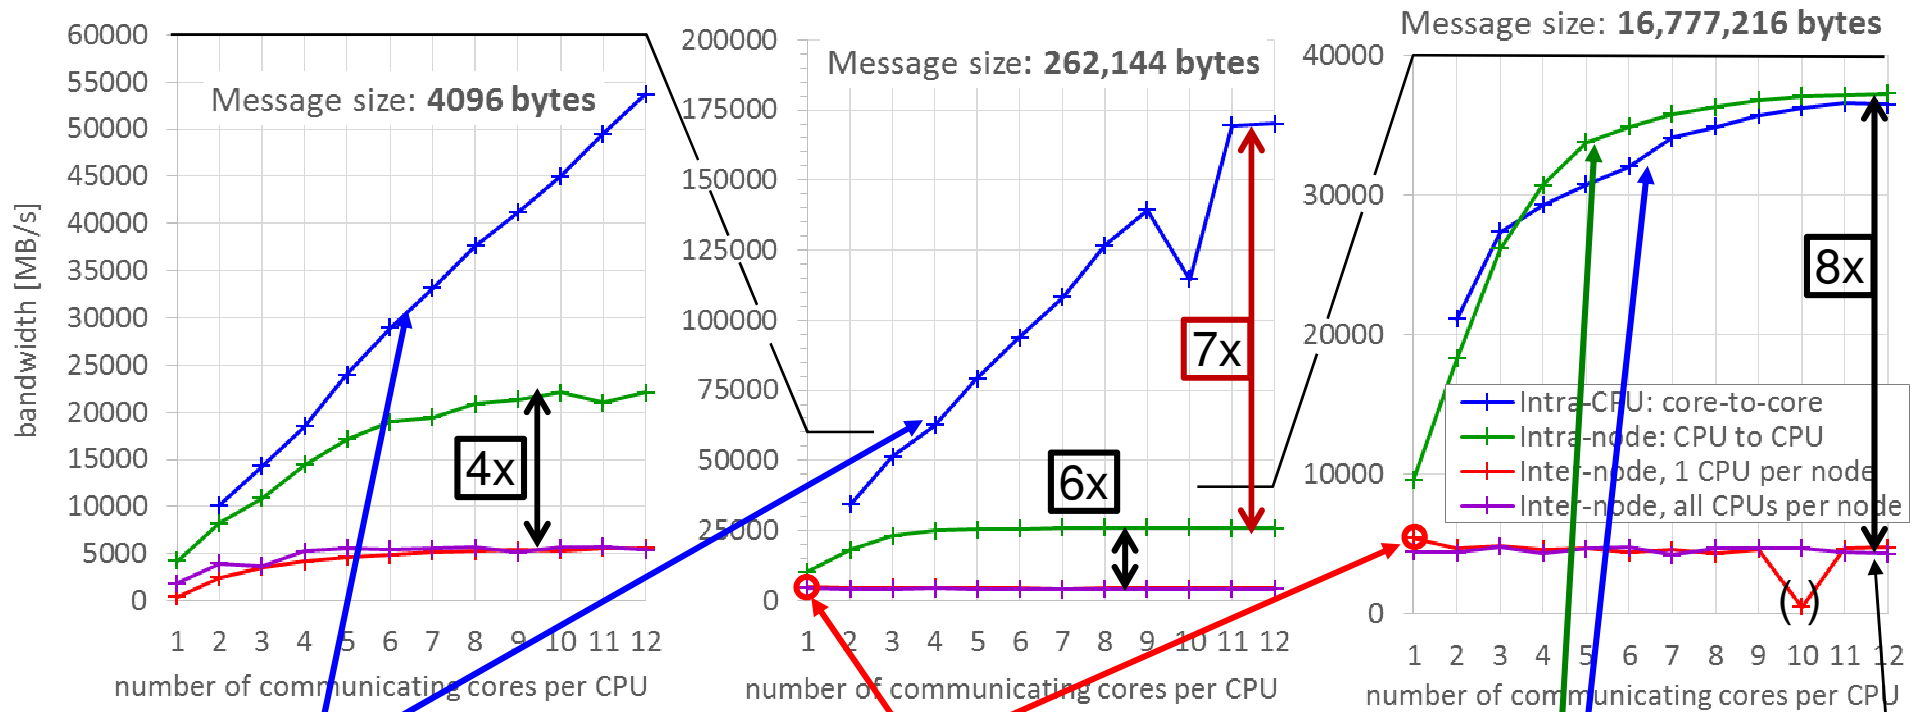
See HLRS online courses <http://www.hlrs.de/training/par-prog-ws/>  
 → Practical → MPI.tar.gz → subdirectory MPI/course/C/1sided/

# Duplex accumulated ring bandwidth per node

(each message is counted twice, as outgoing and incoming)



# Duplex accumulated ring bandwidth per node – scaling vs. asymptotic behavior



Core-to-core:  
Linear scaling for small to medium size messages due to caches

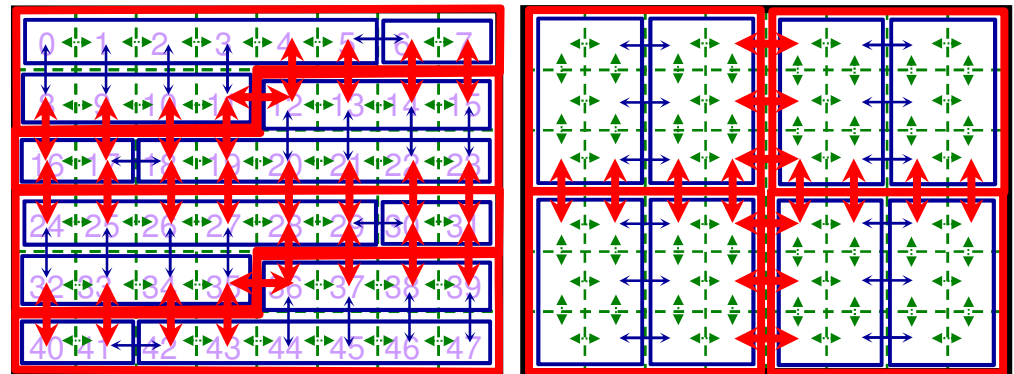
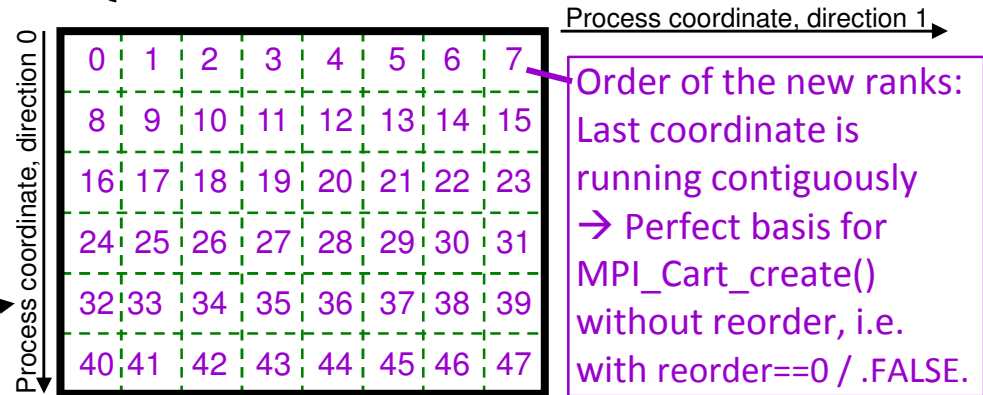
Node-to-node:  
One duplex link by **one core** already fully saturates the network

Core-to-core & CPU-to-CPU:  
**Long messages:**  
Same asymptotic limit through **memory bandwidth**

Result: The limit of accumulated **intra-CPU** and **intra-node** bandwidth is **8x larger** than the limit of accumulated **node-to-node** bandwidth

# Re-numbering on a cluster of SMPs (cores / CPUs / nodes)

- Example with 48 cores on:
  - 4 ccNUMA nodes
  - each node with 2 CPUs ,
  - each CPU with 6 cores
- 2-dim application with 6000 x 8080 gridpoints
  - Minimal communication with 2-dim domain composition with 1000 x 1010 gridpoints/core (shape as quadratic as possible → minimal circumference → minimal halo communication)
  - virtual 2-dim process grid: 6 x 8
- How to locate the MPI processes on the hardware?
  - Using sequential ranks in MPI\_COMM\_WORLD
  - Optimized placement → Proposed algorithm in slides 7-15



**Non-optimal communications:**

- ↔ 26 node-to-node (outer)
- ↔ 20 CPU-to-CPU (middle)
- ↔ 36 core-to-core (inner)

**Optimized placement:**

- ↔ Only 14 node-to-node
- ↔ Only 12 CPU-to-CPU
- ↔ 56 core-to-core

# Hierarchical Cartesian Domain Decomposition

Example:  
24 SMP nodes  
X  
32 cores/node

Per node:  
maximal  
 $8+8+8+8+16+16^*)=$   
**48 or 64<sup>\*)</sup>**  
connections  
to neighbor  
nodes  
\*) with cyclic communication

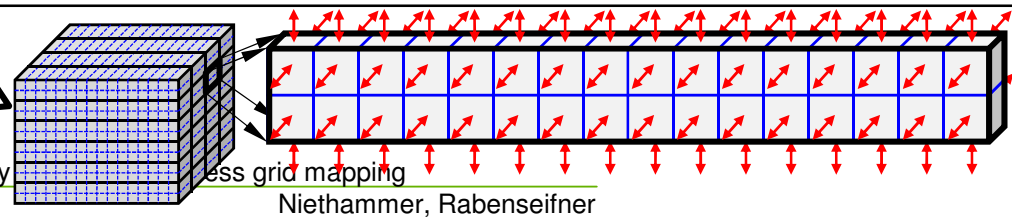
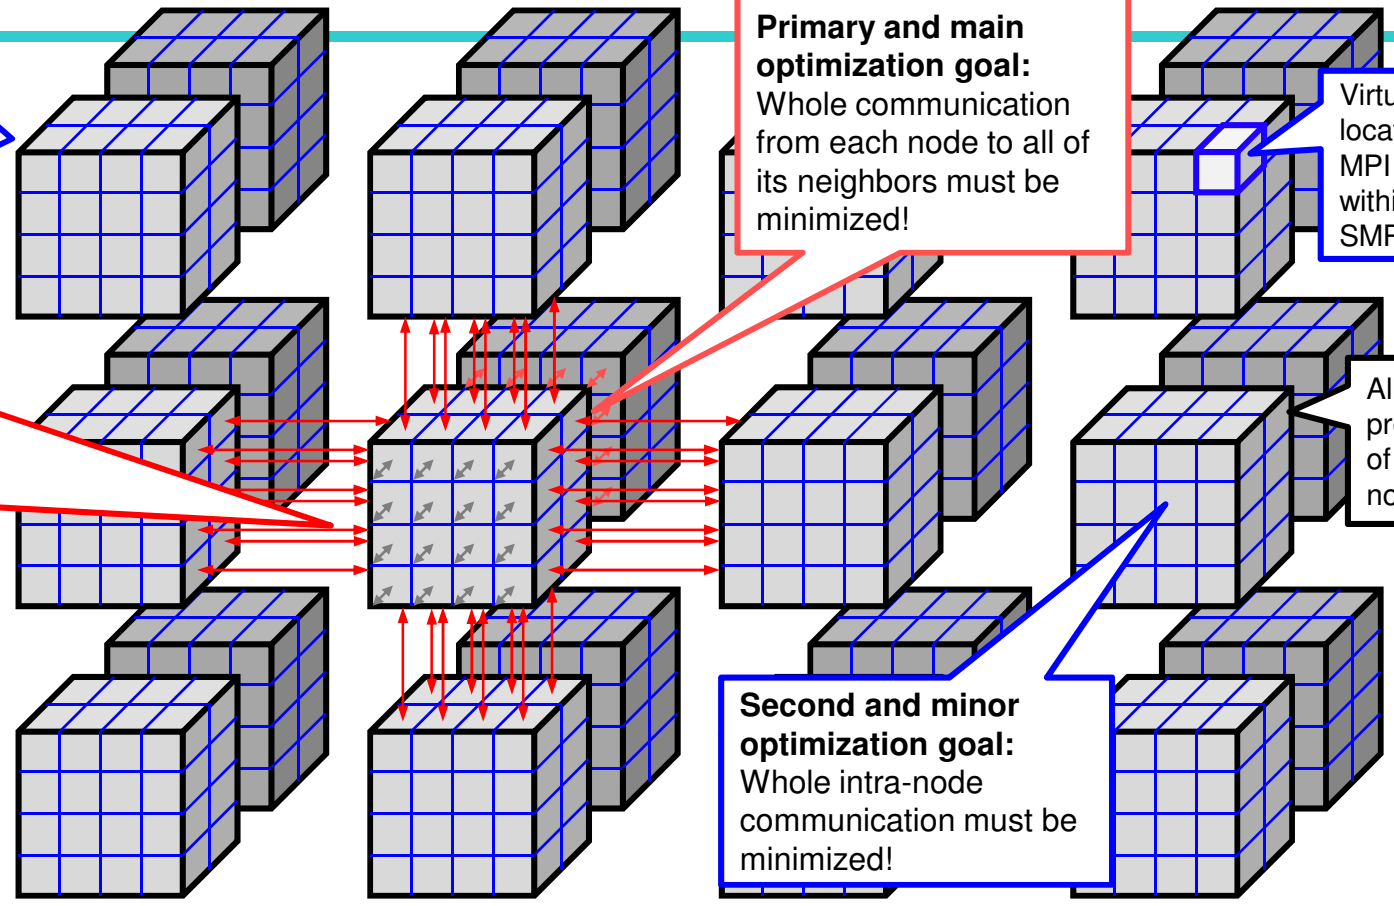
Without  
topology-  
optimization:  
**96** connections  
to other nodes

**Primary and main  
optimization goal:**  
Whole communication  
from each node to all of  
its neighbors must be  
minimized!

Virtual  
location of an  
MPI process  
within an  
SMP node

All MPI  
processes  
of an SMP  
node

**Second and minor  
optimization goal:**  
Whole intra-node  
communication must be  
minimized!



2 or 1.6<sup>\*)</sup> times slower  
communication

Topology  
Slide 8

Niethammer, Rabenseifner

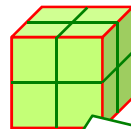
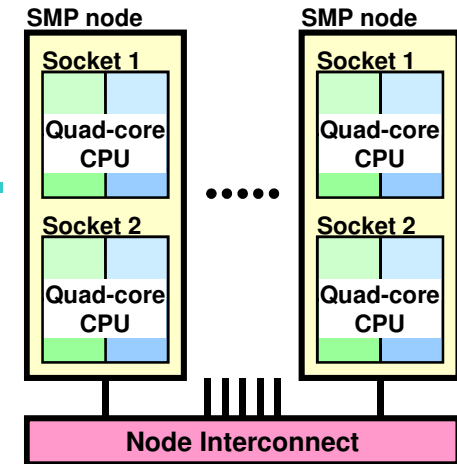
H L R S





# Levels of communication & data access

- Three levels:
  - Between the SMP nodes
  - Between the sockets inside of a ccNUMA SMP node
  - Between the cores of a socket
- On all levels, the communication should be minimized:
  - With 3-dimensional sub-domains:
    - They should be as cubic as possible = minimal surface = minimal communication



Outer surface corresponds to the data communicated to the neighbor nodes in all 6 directions



**Major optimization goal**

Inner surfaces correspond to the data communicated or accessed between the cores inside of a node



Least important

- “as cubic as possible” may be qualified due to different communication bandwidth in each direction caused by sending (fast) non-strided or (slow) strided data

# Back to the problems

1. All MPI libraries provide the necessary interfaces 😊 😊 😊,  
but **without** re-numbering in nearly all MPI-libraries 😞 😞 😞
  - You may substitute `MPI_Cart_create()`  
by the software solution of Bill Gropp (see Bill Gropp, EuroMPI 2018)
2. The existing MPI-3.1 interfaces are not optimal
  - for cluster of ccNUMA node hardware,
    - We substitute `MPI_Dims_create() + MPI_Cart_create()`  
by `MPIX_Cart_weighted_create(... MPIX_WEIGHTS_EQUAL ...)`
  - nor for application specific grid sizes  
or direction-dependent bandwidth
    - by `MPIX_Cart_weighted_create( ... weights ....)`
3. Caution: The application must be prepared for rank re-numbering
  - All communication through the newly created  
**Cartesian communicator with re-numbered ranks!**
  - One must not load data based on `MPI_COMM_WORLD` ranks!



# The new interfaces

## Substitute for / enhancement to existing MPI-1

- MPI\_Dims\_create (size\_of\_comm\_old, ndims, *dims[ndims]* );
- MPI\_Cart\_create (comm\_old, ndims, dims[ndims], periods, reorder, *\*comm\_cart*);

## New:

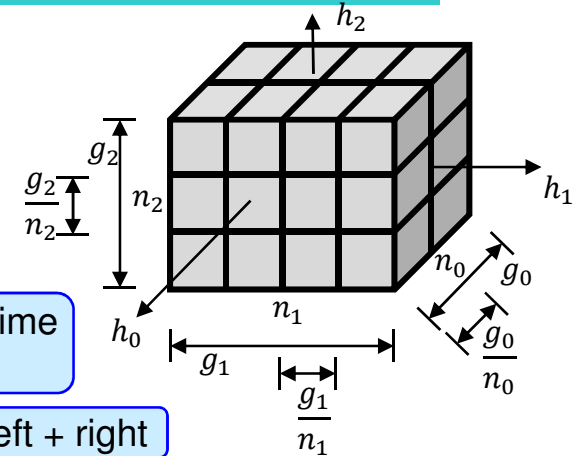
- **MPI\_Cart\_weighted\_create** (  
    /\*IN\*/ MPI\_Comm comm\_old,  
    /\*IN\*/ int ndims,  
    /\*IN\*/ double dim\_weights[ndims], /\*or MPIX\_WEIGHTS\_EQUAL\*/  
    /\*IN\*/ int periods[ndims],  
    /\*IN\*/ MPI\_Info info, /\* for future use, currently MPI\_INFO\_NULL \*/  
    /\*INOUT\*/ int *dims[ndims]*,  
    /\*OUT\*/ MPI\_Comm *\*comm\_cart* );
  - Arguments have same meaning as in MPI\_Dims\_create & MPI\_Cart\_create
  - See next slide for meaning of dim\_weights[ndims]



# The weights $w_i$

## Given:

- $d$ -dimensional Cartesian grid with a total grid size of  $G = \prod_{i=0}^{d-1} g_i$  elements
- The communication cost in each direction  $i = 0, d-1$  is multiplied
  - with a halo width  $h_i$ ,
  - and a communication cost factor  $c_i$ ,



2 = left + right

- total communication cost is  $2g_1g_2h_0c_0 + 2g_0g_2h_1c_1 + 2g_0g_1h_2c_2$
- The weight  $w_i$  is defined as total cost for the communication in one direction:

-  $w_i = 2 \frac{G}{g_i} h_i c_i$  common factors, like  $2G$  or absolute values of  $c_i$ , are not relevant

- With a domain decomposition (i.e., factorization) to  $N = \prod_{i=0}^{d-1} n_i$  nodes, the total communication costs per node is

$$2 \frac{g_1g_2}{n_1n_2} h_0c_0 + 2 \frac{g_0g_2}{n_0n_2} h_1c_1 + 2 \frac{g_0g_1}{n_0n_1} h_2c_2 = \sum_{i=0}^{d-1} \frac{n_i}{N} w_i$$

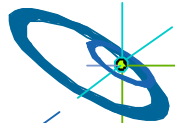
Primarily for the node decomposition and secondarily on core level

- The topology functions have to find a factorization with minimal  $\sum_{i=0}^{d-1} n_i w_i$

→ common factors  $2G$  and  $\frac{1}{N}$  are not relevant →

User level

MPI library level



# The weights $w_i$

## Note that

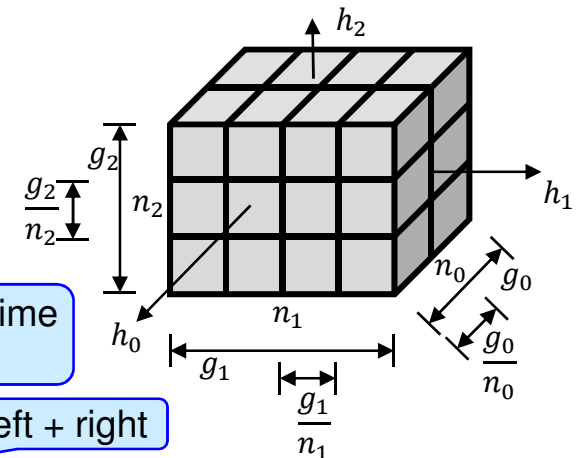
- The proposal for the MPI standard does not discuss the implementation.
- It uses the same figure, but  $n_i$  expresses number of processes, i.e. `dims[i]`, and not number of nodes in one direction, as on previous slide.
- Can be implemented with `MPI_Cart_ml_create_from_types()` → next slide

## Given:

- $d$ -dimensional Cartesian grid with a total grid size of  $G = \prod_{i=0}^{d-1} g_i$  elements
- The communication cost in each direction  $i = 0, d-1$  is multiplied
  - with a halo width  $h_i$ ,
  - and a communication cost factor  $c_i$ ,
- total communication cost is  $2g_1g_2h_0c_0 + 2g_0g_2h_1c_1 + 2g_0g_1h_2c_2$
- The weight  $w_i$  is defined as total cost for the communication in one direction:

$$w_i = 2 \frac{G}{g_i} h_i c_i$$

common factors, like  $2G$  or absolute values of  $c_i$ , are not relevant



User level



# Further Interfaces (1)

If the application wants to choose the hardware levels, then the most simple interface is to choose a split type for each level (except for the last one):

Substitute for / enhancement to existing MPI-1  
 MPI\_Dims\_create ( size\_of\_comm\_old, ndims, *dims* );  
 MPI\_Cart\_create ( comm\_old, ndims, dims, periods, reorder, *\*comm\_cart* );

```
MPI_Cart_ml_create_from_types(MPI_Comm comm_old,
    int nsplit_types,
    int ndims,
    int periods[ndims],
    int split_types[nsplit_types],
    double dim_weights[ndims],
    MPI_Info info,
    /*OUT*/ int dims[ndims], MPI_Comm *comm_cart );
```

e.g., with  
 25\*25\*24 = **15000 processes**  
 on **625 ccNUMA nodes** with  
**2 CPUs/node** and **12 cores/CPU**

e.g., nsplit\_types=2, split\_types=  
 { MPI\_COMM\_TYPE\_SHARED,  
 OMPI\_COMM\_TYPE\_NUMA  
 within OpenMPI, or for further  
 splitting: MPIX\_COMM\_TYPE\_  
 HW\_TOPOLOGY }

e.g., 3 dimensions with a data  
 grid with 1000 x 1100 x 950  
 elements → dim\_weights[] =  
 { 1.0/1000, 1.0/1100, 1.0/950 }

Rank mapping is based on:

- Node level: **625** = 5 x 25 x 5
- CPU level: **2** = 2 x 1 x 1
- Core level: **12** = 3 x 1 x 4

Result (product): **30 x 25 x 20**

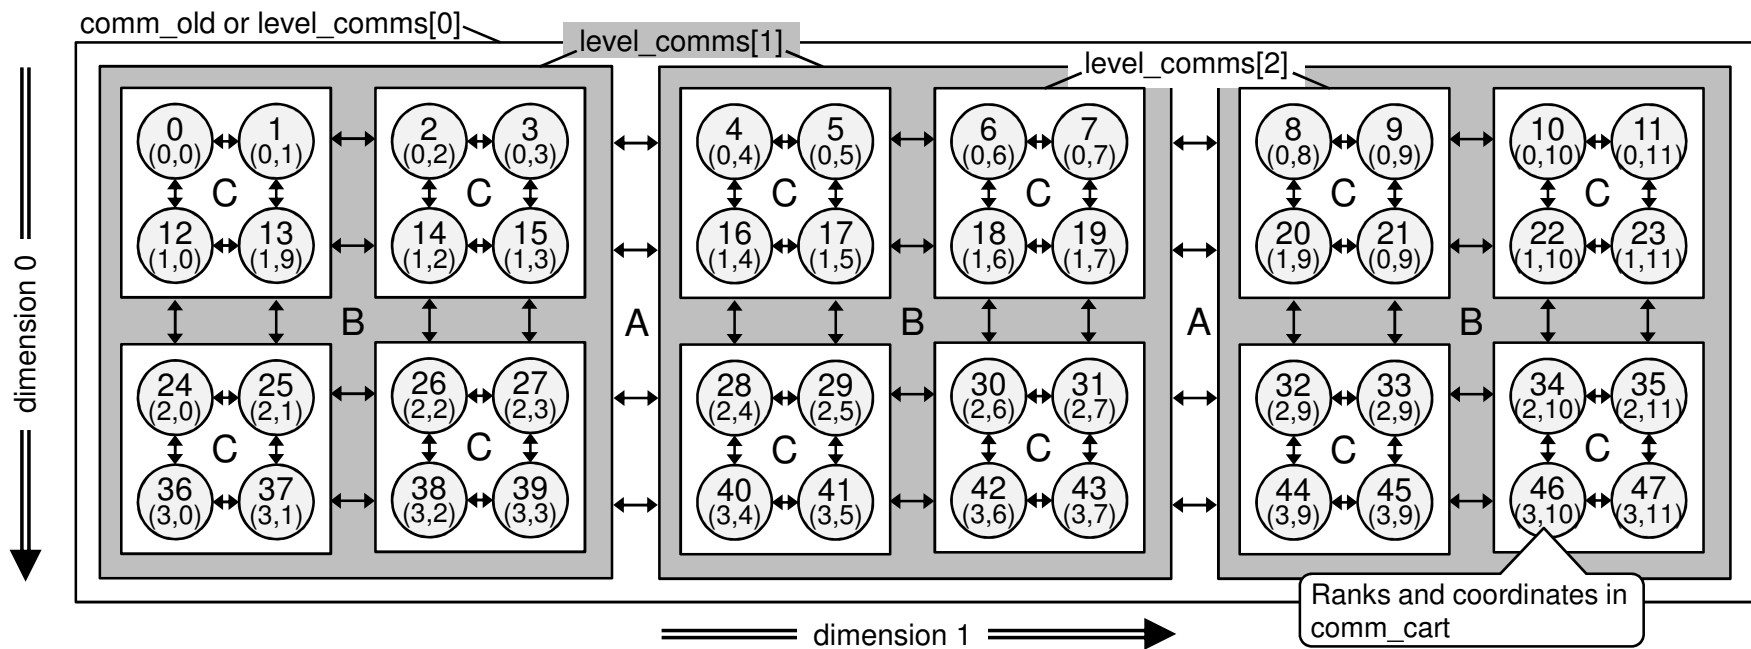
The Cartesian communicator reflects this result: **30 x 25 x 20**

Note that MPI\_Dims\_create() would produce  
**25 x 25 x 24**  
 which would never fit to the needed node-level distribution  
**5 x 25 x 5**

Next steps for the application:  
**MPI\_Comm\_rank ( comm\_cart, &my\_rank );**  
**MPI\_Cart\_coords ( comm\_cart, my\_rank, ndims, *coords* )**

## 2-D Cartesian Example (proposed for the MPI standard)

- with 48 processes and ndims=2,
- dim\_weights=(1.0/4, 1.0/12), and
- nsplit\_types=2 (e.g. splitting into ccNUMA nodes and those into CPUs)
- or appropriate level\_comms[0] .. [2]



The optimization chooses an appropriate re-ordering of the ranks of `comm_cart` that  
 1<sup>st</sup> communication A (e.g., node to node) is minimized,  
 2<sup>nd</sup> B (e.g., CPU to CPU) is minimized, and 3<sup>rd</sup> C (e.g., core to core) is minimized

## Further Interfaces (2)

If the application wants to choose the hardware levels, but appropriate split types (as for `MPI_COMM_SPLIT_TYPE()`) do not exist, then the splitting can be done by the application and an array with the hierarchical set of communicators is the input:

Substitute for / enhancement to existing `MPI-1`

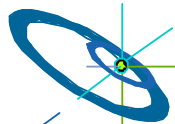
```
MPI_Dims_create ( size_of_comm_old, ndims, dims );  
MPI_Cart_create ( comm_old, ndims, dims, periods, reorder, *comm_cart );
```

```
MPI_Cart_ml_create_from_comms(int nlevels,  
MPI_Comm level_comms[nlevels],  
int ndims, double dim_weights[ndims], int periods[ndims], MPI_Info info,  
/*OUT*/ int dims[ndims], MPI_Comm *comm_cart );
```

e.g., `nlevels=3`, and `level_comms[0]` is `comm_old`, `level_comms[1 and 2]` are the result recursively called `MPI_Comm_split_type` with the `type_levels` from previous slide.

Same as above

Same as above





## Further Interfaces (3) – the basis

**MPI\_Cart\_ml\_create\_from\_comms**(... **level\_comms**, ...)  
can be implemented based on **MPI\_Dims\_ml\_create**(),  
which can be implemented through **MPI\_Dims\_weighted\_create**()  
provided that the **level\_comms** have same size on each level

**MPI\_Dims\_weighted\_create** ( int nnodes, int ndims, double dim\_weights[ndims],  
int periods[ndims], MPI\_Info info, */\*INOUT\*/ int dims[ndims]*);

**MPI\_Dims\_ml\_create** ( int nnodes, int ndims, double dim\_weights[ndims],  
int nlevels, int sizes[nlevels], int periods[ndims], MPI\_Info info,  
*/\*OUT\*/ int dims\_ml[ndims][nlevels]*);

Multi-level  
info



## Further Interfaces (3) – the basis (a)

```
MPI_Dims_weighted_create ( int nnodes, int ndims, double dim_weights[ndims],  
                           int periods[ndims], MPI_Info info, /*INOUT*/ int dims[ndims]);
```

The factorization of nnodes into the array dims is chosen in a way to minimize the communication according to the given weights. Note that this means that this function looks for a factorization with minimal sum  $\sum_i \text{dims}[i] \cdot \text{dims\_weight}[i]$ .

Rationale. As shown in Figure 7.3, the total communication cost per node (left and right, in all dimensions) would be  $C = 2 \sum_i w_i / \prod_{j, j \neq i} d_j$  and with  $\text{nnodes} = N = \prod_j d_j$ , the cost is  $C = \frac{2}{N} \sum_i w_i d_i$  (End of rationale.)

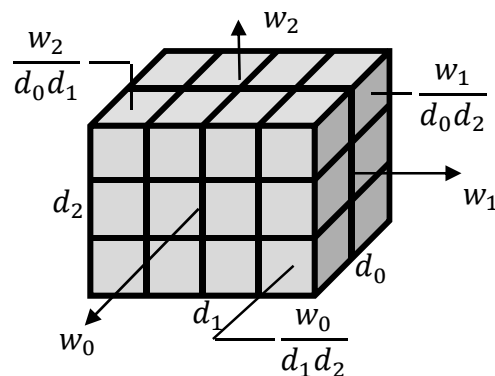


Figure 7.3: Communication model:  
The  $\text{dim\_weights}[i] = w_i$  express the total communication cost in the direction of dimension  $i$ . The communication cost of one process in direction  $i$  is therefore  $w_i / \prod_{j, j \neq i} d_j$  with  $\text{dims}[i] = d_i$ .

## Further Interfaces (3) – the basis (b)

```
MPI_Dims_ml_create ( int nnodes, int ndims, double dim_weights[ndims],  
                    int nlevels, int sizes[nlevels], int periods[ndims], MPI_Info info,  
                    /*OUT*/ int *dims_ml[ndims][nlevels] );
```

It uses internally `MPI_Dims_weighted_create()` for each level → next slide

int \*dims\_ml

- represents in C an C++ a contiguous two dimensional array `dims_ml[ndims][nlevels]`
- and `dims_ml(ndims,nlevels)` in Fortran



# Further Interfaces

## Summary of (1-3)

e.g., with  
 $25 \times 25 \times 24 = 15000$  processes  
 on **625 ccNUMA nodes** with  
**2 CPUs/node** and **12 cores/CPU**

e.g., {  
 MPI\_COMM\_TYPE\_SHARED,  
 OMPI\_COMM\_TYPE\_NUMA  
 within OpenMPI, or for further  
 splitting: MPIX\_COMM\_TYPE\_  
 HW\_TOPOLOGY}

e.g., 3 dimensions with a data  
 grid with 1000 x 1100 x 950  
 elements → dim\_weights[] =  
 { 1.0/1000, 1.0/1100, 1.0/950 }

```
MPI_Cart_ml_create_from_types(MPI_Comm comm_old,
                              int nsplit_types,
                              int ndims,
                              int periods[ndims],
                              int split_types[nsplit_types],
                              double dim_weights[ndims],
                              MPI_Info info,
                              /*OUT*/ int dims[ndims], MPI_Comm *comm_cart);
```

Rank mapping is based on:

- Node level: **625** = 5 x 25 x 5
- CPU level: **2** = 2 x 1 x 1
- Core level: **12** = 3 x 1 x 4

Result (product): **30 x 25 x 20**

The Cartesian communicator reflects this result: **30 x 25 x 20**

Next steps:  
**MPI\_Comm\_rank ( comm\_cart, &my\_rank );**  
**MPI\_Cart\_coords ( comm\_cart, my\_rank, ndims, coords)**

```
MPI_Cart_ml_create_from_comms(int nlevels,
                              MPI_Comm level_comms[nlevels],
                              int ndims,
                              double dim_weights[ndims],
                              int periods[ndims],
                              MPI_Info info,
                              /*OUT*/ int dims[ndims], MPI_Comm *comm_cart);
```

e.g., level\_comms[0] is comm\_old, level\_comms[1 and 2] are the result recursively called MPI\_Comm\_split\_type with the type\_levels from above.

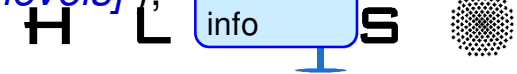
Same as above

Same as above

```
MPI_Dims_weighted_create ( int nnodes, int ndims, double dim_weights[ndims],
                          int periods[ndims], MPI_Info info, /*INOUT*/ int dims[ndims]);
```

```
MPI_Dims_ml_create ( int nnodes, int ndims, double dim_weights[ndims],
                    int nlevels, int sizes[nlevels], int periods[ndims], MPI_Info info,
                    /*OUT*/ int dims_ml[ndims][nlevels]);
```

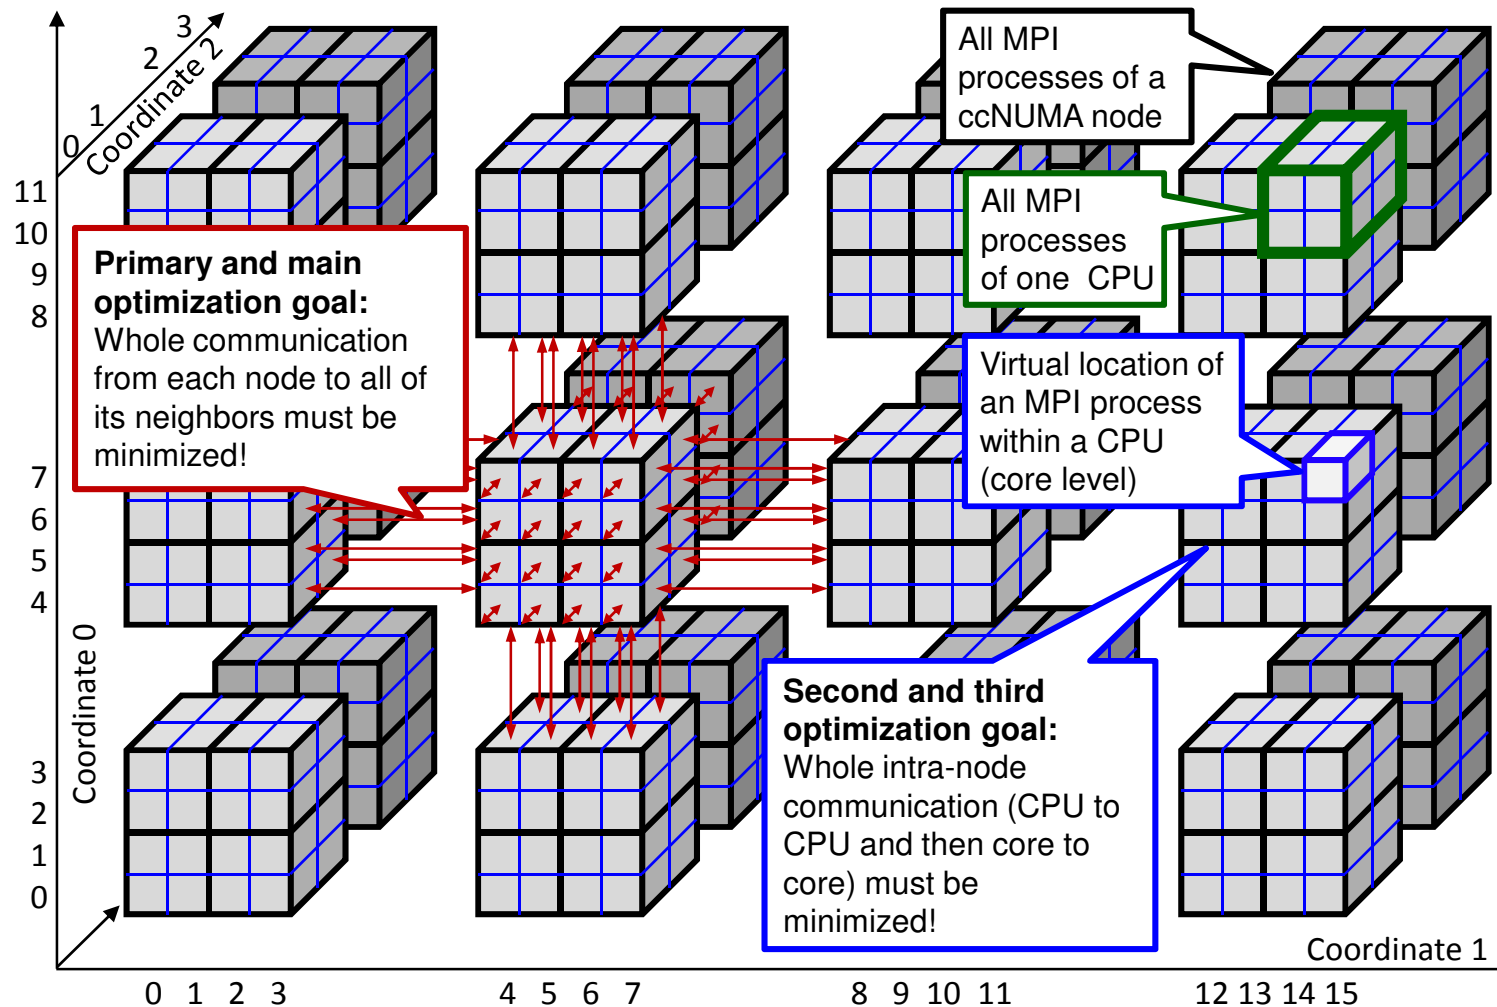
Multi-level  
 info



Substitute for / enhancement to existing MPI-1  
 MPI\_Dims\_create ( size\_of\_comm\_old, ndims, dims );  
 MPI\_Cart\_create ( comm\_old, ndims, dims, periods,  
 reorder, \*comm\_cart );

# Hierarchical Cartesian Domain Decomposition

(figure used in the proposal for the MPI standard)



## Further Interfaces (4)


From previous slides

```
MPI_Dims_ml_create ( int nnodes, int ndims, double dim_weights[ndims],  
                    int nlevels, int sizes[nlevels], int periods[ndims], MPI_Info info,  
                    /*OUT*/ int dims_ml[ndims][nlevels] );
```

```
MPI_Cart_ml_create (MPI_Comm comm_old, int ndims, int *periods,  
                   int nlevels, int dims_ml[ndims][nlevels], MPI_Info info,  
                   /*OUT*/ int *dims, MPI_Comm *comm_cart );
```

This interface requires that comm\_old is **ranked sequentially in the hardware**

We proposed the algorithm in

- Christoph Niethammer and Rolf Rabenseifner. 2018. Topology aware Cartesian grid mapping with MPI. EuroMPI 2018. <https://eurompi2018.bsc.es/>  
→ Program → Poster Session → Abstract+Poster
- <https://fs.hlr.de/projects/par/mpi/EuroMPI2018-Cartesian/>  
→ All info + slides  + software
- <http://www.hlr.de/training/par-prog-ws/>  
→ Practical → MPI.tar.gz → MPI/course/C/eurompi18/

Here, you get the  
new **optimized  
interface** +  
**implementation +  
documentation**

MPIX\_Dims\_weighted\_create() is based on the ideas in:

- Jesper Larsson Träff and Felix Donatus Lübbe. 2015. Specification Guideline Violations by MPI Dims Create. In *Proceedings of the 22nd European MPI Users' Group Meeting (EuroMPI '15)*. ACM, New York, NY, USA, Article 19, 2 pages.

Substitute for / enhancement to existing MPI-1

MPI\_Dims\_create ( size\_of\_comm\_old, ndims, *dims* );

MPI\_Cart\_create ( comm\_old, ndims, dims, periods,  
reorder, *\*comm\_cart* );

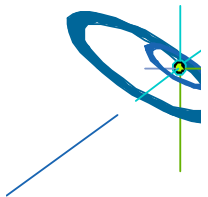


## Implementation Remarks

- The portable MPIX routines internally use `MPI_Comm_split_type(..., MPI_COMM_TYPE_SHARED, ...)` to split `comm_old` into ccNUMA nodes,
- plus (may be) additional splitting into NUMA domains.
- With using hyperthreads, it *may be helpful* to apply **sequential** ranking to the hyperthreads,
  - i.e., in `MPI_COMM_WORLD`, ranks 0+1 should be
    - **the first two hyperthreads**
    - of the first core
    - of the first CPU
    - of the first ccNUMA node
- Especially with weights  $w_i$  based on  $\frac{G}{g_i}$ , it is important
  - that the data of the grid points is **not** read in based on (**old**) ranks in `MPI_COMM_WORLD`,
  - because the domain decomposition must be done based on **`comm_cart`** and its dimensions and (**new**) ranks

# Internal implementation plan

- **MPIX\_Cart\_weighted\_create(...)**
  - chooses available and useful types for splitting, e.g., {MPI\_COMM\_TYPE\_SHARED, OMPI\_COMM\_TYPE\_NUMA or MPIX\_COMM\_TYPE\_HALFNODE}
  - **MPIX\_Cart\_ml\_create\_from\_types(...)**
- **MPIX\_Cart\_ml\_create\_from\_types(...)**
  - loop over MPIX\_Comm\_split\_type
  - **MPIX\_Cart\_ml\_create\_from\_comms(...)**
- **MPIX\_Cart\_ml\_create\_from\_comms(...)**
  - must calculate level\_sizes[nlevels] and whether they are equally sized within the same level
  - if (equally-sized) then
    - **MPIX\_Dims\_ml\_create(...)**
      - Appropriate renumbering based on dims\_ml and the level\_comms
      - Calculation of dims[] & creation of comm\_cart → **MPI\_Cart\_create(...)** without reorder
  - else, e.g., algorithm of Thorsten Hoefler
- **MPIX\_Cart\_ml\_create(...)** → Usable only for sequentially ranked comm\_old
  - Appropriate renumbering based on dims\_ml and the sequential comm\_old
  - Calculation of dims[] & creation of comm\_cart → **MPI\_Cart\_create(...)** without reorder
- **MPIX\_Dims\_ml\_create(...)**
  - **MPIX\_Dims\_weighted\_create(...)** on each level
- **MPIX\_Dims\_weighted\_create(...)**
  - This is the important new base routine with a new fast brute force algorithm





# Benchmark: halo\_irecv\_send\_toggle\_3dim\_grid\_solution.c

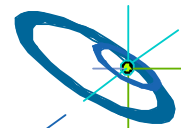
- Input per measurement, e.g. on 8 nodes x 2 CPUs x 12 cores: Example
  - cart\_method: 2
    - 0=end, 1=Dims\_create+Cart\_create,
    - 2=Cart\_weighted\_create(MPIX\_WEIGHTS\_EQUAL),
    - 3=dito(weights), 4=dito manually, 5=Cart\_ml\_create(dims\_ml)
  - start grid sizes integer start values 1 2 4
    - Using MPI\_Type\_vector, for each dimension a pair of blocklength&stride 0 0 0 0 0 0
      - weights (double values) (only with cart\_method==4) 1.00 0.50 0.25
      - number of hardware levels (only with cart\_method==5) 3
    - dims\_ml: for each of the 3 Cartesian dimensions a list of 3 dimensions from outer to inner hardware level, e.g., 8 nodes x 2 CPUs x 12 cores are split into 1x2x4 nodes x 2x1x1 CPUs x 2x3x2 cores
      - dims\_ml[d=0] = 1 2 2
      - dims\_ml[d=1] = 2 1 3
      - dims\_ml[d=2] = 4 1 2
- Input can be concatenated to one line per experiment:
  - 1 1 2 4 0 0 0 0 0 0
  - 2 1 2 4 0 0 0 0 0 0
  - 3 1 2 4 0 0 0 0 0 0
  - 4 1 2 4 0 0 0 0 0 0 4. 2. 1.
  - 5 1 2 4 0 0 0 0 0 0 3 1 2 2 2 1 3 4 1 2
  - 3 2 2 2 256 1024 4 32 0 0
  - 0

0 0 = contiguous

← 3





Start a 16-node batch-job with your own input file:  
**Report your acceleration factors to the course group**

examples for strided data in direction 0 & 1



## Additional Remarks

- Caution with stdout and stdin when switching I/O from process `world_rank==0` to `cart_rank==0`:
  - **Before** establishing the new `comm_cart`, all I/O on stdout/stdin is done by `world_rank==0` (in **`MPI_COMM_WORLD`**)
  - **After** establishing the new `comm_cart`, all I/O on stdout/stdin is done by `cart_rank==0` (in `comm_cart`)
  - In between, we recommended (although it is not guaranteed that an *output on comm\_cart* may overtake an *output on MPI\_COMM\_WORLD*):
    - **`MPI_Barrier(MPI_COMM_WORLD);`**
    - **`sleep(1);`** // costs nearly nothing, e.g., 30 Mio € TCO/year / (365 days/year \* 24 hours/day \* 3600 sec/hour) \* 1 sec = 1€
    - **`MPI_Barrier(comm_cart);`**
- The following slide shows the win through the re-ranking by the new routines:
  - Less % is better – the communication time reduction factors are:

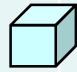
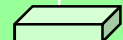
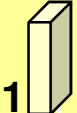
- 1.1-1.2 
- 1.75 
- 2.75 
- 4.5-5.0 

Topology aware MPI process grid mapping

Slide 26

Niethammer, Rabenseifner



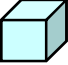
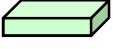
Halosize/process ≈ 26 MB <small>Depend on chosen dims</small>		MPI_Dims_create + MPI_Cart_create		MPIX_Cart_weighted_create(MPIX_WEIGHTS_EQUAL)		MPIX_Cart_weighted_create(...weights...)	
Base grid sizes	Nodes x CPUs x cores	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communication time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]
<b>2 x 2 x 2</b>  <small>Base factors, not absolute values</small>	<b>8</b> x2x12	84.545 = baseline	8 = <b>8</b> x 1 x 1 6 = <b>1</b> x 2 x 3 4 = <b>1</b> x 1 x 4	52.666 = 62%	8 = <b>2</b> x 2 x 2 6 = <b>2</b> x 1 x 3 4 = <b>2</b> x 1 x 2	48.556 = 57%	8 = <b>2</b> x 2 x 2 6 = <b>2</b> x 1 x 3 4 = <b>2</b> x 1 x 2
	<b>64</b> x2x12	194.856 = baseline	16= <b>16</b> x 1 x 1 12= <b>4</b> x 1 x 3 8= 1 x 2 x 4	73.756 = 38%	16= <b>4</b> x 2 x 2 12= <b>4</b> x 1 x 3 8= <b>4</b> x 1 x 2	72.051 = 37%	16= <b>4</b> x 2 x 2 12= <b>4</b> x 1 x 3 8= <b>4</b> x 1 x 2
	<b>512</b> x2x12	247.631 = baseline	32= <b>32</b> x 1 x 1 24= <b>16</b> x 1x1.5 16= 1 x 2 x 8	85.530 = 35%	32= <b>8</b> x 2 x 2 24= <b>8</b> x 1 x 3 16= <b>8</b> x 1 x 2	85.491 = 35%	32= <b>8</b> x 2 x 2 24= <b>8</b> x 1 x 3 16= <b>8</b> x 1 x 2
<b>1 x 2 x 4</b>  <small>By default, communication time strongly depends on cuboid's direction</small>	<b>8</b> x2x12	172.850 = baseline	8 = <b>8</b> x 1 x 1 6 = <b>1</b> x 2 x 3 4 = 1 x 1 x 4	63.796 = 37%	8 = <b>2</b> x 2 x 2 6 = <b>2</b> x 1 x 3 4 = <b>2</b> x 1 x 2	37.953 = 22%	4 = <b>1</b> x 2 x 2 6 = <b>2</b> x 1 x 3 8 = <b>4</b> x 1 x 2
	<b>64</b> x2x12	360.364 = baseline	16= <b>16</b> x 1 x 1 12= <b>4</b> x 1 x 3 8= 1 x 2 x 4	91.524 = 25%	16= <b>4</b> x 2 x 2 12= <b>4</b> x 1 x 3 8= <b>4</b> x 1 x 2	74.199 = 21%	8 = <b>2</b> x 2 x 2 12= <b>4</b> x 1 x 3 16= <b>8</b> x 1 x 2
	<b>512</b> x2x12	457.858 = baseline	32= <b>32</b> x 1 x 1 24= <b>16</b> x 1x1.5 16= 1 x 2 x 8	125.468 = 27%	32= <b>8</b> x 2 x 2 24= <b>8</b> x 1 x 3 16= <b>8</b> x 1 x 2	93.615 = 20%	16= <b>4</b> x 2 x 2 24= <b>8</b> x 1 x 3 32= <b>16</b> x 1 x 2
<b>4 x 2 x 1</b>  <small>On Cray XC40 Hazel hen at HLRS Stuttgart, Jan 2019</small>	<b>8</b> x2x12	40.050 = baseline	8 = <b>8</b> x 1 x 1 6 = <b>1</b> x 2 x 3 4 = <b>1</b> x 1 x 4	59.421 =148%	8 = <b>2</b> x 2 x 2 6 = <b>2</b> x 1 x 3 4 = <b>2</b> x 1 x 2	36.778 =92%	16 = <b>4</b> x 2 x 2 6 = <b>2</b> x 1 x 3 2 = <b>1</b> x 1 x 2
	<b>64</b> x2x12	78.503 = baseline	16= <b>16</b> x 1 x 1 12= <b>4</b> x 1 x 3 8= <b>1</b> x 2 x 4	100.203 =128%	16= <b>4</b> x 2 x 2 12= <b>4</b> x 1 x 3 8= <b>4</b> x 1 x 2	69.802 =89%	32= <b>8</b> x 2 x 2 12= <b>4</b> x 1 x 3 4 = <b>2</b> x 1 x 2
	<b>512</b> x2x12	103.002 = baseline	32= <b>32</b> x 1 x 1 24= <b>16</b> x 1x1.5 16= <b>1</b> x 2 x 8	93.189 = 90%	32= <b>8</b> x 2 x 2 24= <b>8</b> x 1 x 3 16= <b>8</b> x 1 x 2	85.044 =83%	64= <b>16</b> x 2 x 2 24= <b>8</b> x 1 x 3 8= <b>4</b> x 1 x 2

Used process dimensions

Internally applied dimensions on each hardware level

Optimized communication times are nearly independent of cuboid's form

## As Exercise: To do (1)

- `cp ~/MPI/course/C/Ch9/MPIX_*/*` .
  - You get the benchmark skeleton `halo_irecv_send_toggle_3dim_grid_skel.c`
  - And all `MPIX_*.c` files and the header `MPIX_interface_proposal.h`
- `mpicc -o halo_skel.exe halo_irecv_send_toggle_3dim_grid_skel.c MPIX_*.c`
- First test with non-optimized `cart_method==1`, i.e., `MPI_Dims_create + MPICart_create`
  - Choose your batch job: `halo_skel_[LRZ|VSC|HLRS].sh` which contains
    - Number of nodes and cores/node
    - and, e.g.,  
`mpirun -np 192 ./halo_skel.exe < input-skel.txt`
  - Start your batchjob
  - Try to understand the output:
    - It contains two experiments: a grid with cubic  and one with non-cubic  ratio
    - The number of MPI processes, e.g. 192, is factorized  
→ domain decomposition into, e.g., 8 x 6 x 4 processes
    - The measurements are done for 10 global gridsizes
    - The domain decomposition implies the local gridsizes
    - The local gridsizes imply the size of the halos in each direction
    - → the sum of the time for the communication into the 3 dimensions x 2 directions (left+right)

```
1 2 2 2 0 0 0 0 0 0
1 1 2 4 0 0 0 0 0 0
0
```

See  
next slide

# Exercise: To do (2)

**cart\_method = 1**

start grid sizes integer start values for 3 dimensions = 2 2 2  
 blocklength & sgtride pairs for each of the 3 dimensions = 1 0 1 0 1 0

**Creating the Cartesian communicator** and further input arguments:

cart\_method == 1: **MPI\_Dims\_create** + **MPI\_Cart\_create**  
 [MPI\_Barrier and switching to output via stdout through rank==0 in comm\_cart]

ndims=3 dims= 8 6 4

message size      transfertime      duplex bandwidth per process and neighbor (grid&halo in #floats)

message size	transfertime	duplex bandwidth	gridsizes total=	per process=	halosizes=
128 bytes	34.537 usec	3.706 MB/s	16 12 12	2 2 3	16= 6 + 6 + 4
432 bytes	39.840 usec	10.843 MB/s	24 24 24	3 4 6	54= 24 + 18 + 12
1728 bytes	41.122 usec	42.021 MB/s	48 48 48	6 8 12	216= 96 + 72 + 48
6688 bytes	23.961 usec		96 96 92	12 16 23	836= 368 + 276 + 192
25576 bytes	93.703 usec		184 186 184	23 31 46	3197= 1426 + 1058 + 713
104408 bytes	271.721 usec		376 372 372	47 62 93	13051= 5766 + 4371 + 2914
411192 bytes	1033.001 usec	0.056 MB/s	744 738 740	93 123 185	51399= 22755+17205 + 11439
1636392 bytes	4398.680 usec	372.019 MB/s	1480 1476 1476	185 246 369	204549= 90774+68265 + 45510
6561336 bytes	18173.518 usec	361.038 MB/s	2960 2958 2956	370 493 739	820167=364327+273430+182410
<b>26194104 bytes</b>	<b>76132.216 usec</b>	344.061 MB/s	<b>5912</b> 5910 5908	739 <b>985</b> <b>1477</b>	<b>3274263=1454845+1091503+727915</b>

192 processes:  
Input for  
MPI\_Dims\_create()

These base values (per process) are multiplied with  $\sqrt[3]{\#processes}$  and then with 1, 2, 4, 8, ... 512, e.g.,  $2 \cdot \sqrt[3]{192} \cdot 512 = 5912$  (rounded to a multiple the dimension)

= 2 2 2

divide by dims[i]

First value for our table

5912

**cart\_method = 1**

\* 2 directions \* 4 byte

start grid sizes integer start values for 3 dimensions = 1 2 4  
 blocklength & sgtride pairs for each of the 3 dimensions = 0 0 0 0 0 0

cart\_method == 1: **MPI\_Dims\_create** + **MPI\_Cart\_create**

ndims=3 dims= 8 6 4

message size      transfertime      duplex bandwidth per process and neighbor (grid&halo in #floats)

message size	transfertime	duplex bandwidth	gridsizes total=	per process=	halosizes=
160 bytes	14.720 usec	10.870 MB/s	8 12 24	1 2 6	20= 12 + 6 + 2
...	...	...	...	...	...
<b>34936960</b>	<b>156869.278 usec</b>	222.714 MB/s	2960 5910 11816	370 985 2954	4367120=2909690+1092980+364450

multiply for 2-d halo array size

Same values, because MPI\_Dims\_create() factorizes the #processes independent from the user's gridsizes.

Second value for our table



# Exercise: To do (3)

- **Fill in the table**

Given from MPI\_Dims\_create()

**Nodes**  
CPUs  
cores

d=0: **8** = **8** x 1 x 1  
 d=1: **6** = **1** x 2 x 3  
 d=2: **4** = **1** x 1 x 4  


---

 Total 192 = **8 x 2 x 12**

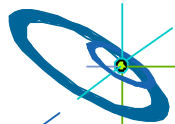
**Please, calculated by hand:**  
 Fill in maximal factors.  
 Factorize first the cores  
 and start with d=2.  
 Then the CPUs & then the nodes.  
 (All based on sequential ranking of MPI\_COMM\_WORLD)

Defined in batch job + hardware knowledge

Execution time of **largest** grid and halo size of both measurements

Given by batchjob and hardware

Halosize/process ~= 26 MB	MPI_Dims_create + MPI_Cart_create	MPI_Comm_create(MPIX_WEIGHTS_EQUAL)	MPIX_Cart_weighted_create(...weights...)
Base grid sizes	Nodes x CPUs x cores Communicat. time [ms] dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communication time [ms] dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communicat. time [ms] dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]
2 x 2 x 2	___ x ___ = baseline ___ = ___ x ___ ___ = ___ x ___ ___ = ___ x ___		
1 x 2 x 4	Same as above = baseline Same as above		



## Exercise: To do (4)

- cp halo\_irecv\_send\_toggle\_3dim\_grid\_skel.c halo\_optim.c
- Edit halo\_optim.c
  - On lines 160, 165, 171, and 184, substitute the `/* TODO: ... */` by correct code

```
153 if (cart_method == 1) {
154     if (my_world_rank==0) printf("cart_method == 1: MPI_Dims_create + MPI_Cart_create\n");
155     MPI_Dims_create(size, ndims, dims);
156     MPI_Cart_create(MPI_COMM_WORLD, ndims, dims, periods, 0, &comm_cart);
157 } else if (cart_method == 2) {
158     if (my_world_rank==0) printf("cart_method == 2: MPIX_Cart_weighted_create( MPIX_WEIGHTS_EQUAL )\n");
160     /* TODO: Appropriate call to MPIX_Cart_weighted_create(...) with MPIX_WEIGHTS_EQUAL
161        instead of calling MPI_Dims_create() and MPI_Cart_create() as in method 1 */
163 } else if (cart_method == 3) {
165     /* TODO: Appropriate calculation of weights[ ] based on gridsize_avg_per_proc_startval[ ] */
167     if (my_world_rank==0) { printf("cart_method == 3: MPIX_Cart_weighted_create( weights := _____TODO_____)\n");
168         printf("weights= "); for (d=0; d<ndims; d++) printf(" %lf",weights[d]); printf("\n");
169     }
171     /* TODO: Appropriate call to MPIX_Cart_weighted_create(...) with weights
172        instead of MPIX_WEIGHTS_EQUAL as in method 2 */
174 } else if (cart_method == 4) {
175     for (d=0; d<ndims; d++) weights[d] = 4.0 / gridsize_avg_per_proc_startval[d];
176     if (my_world_rank==0) { printf("cart_method == 4: MPIX_Cart_weighted_create( manual weights )\n");
177         printf("weights (double values) for %d dimensions (e.g., ", ndims);
178         for (d=0; d<ndims; d++) printf(" %lf",weights[d]); printf(") ?\n");
179         for (d=0; d<ndims; d++) scanf("%lf",&weights[d]);
180         printf("weights= "); for (d=0; d<ndims; d++) printf(" %lf",weights[d]); printf("\n");
181     }
182     MPI_Bcast(weights, ndims, MPI_DOUBLE, 0, MPI_COMM_WORLD);
184     /* TODO: Appropriate call to MPIX_Cart_weighted_create(...)
185        same as in method 3, but without the calculation of the weights */
187 } else { ...
```

# Exercise: To do (5)

- `mpicc -o halo_optim.exe halo_optim.c MPIX_*.c`
- Check: `diff halo_optim.c halo_irecv_send_toggle_3dim_grid_solution.c`
- Now, use all three `cart_method==1, 2, 3`
  - Choose your batch job:
    - `halo_optim_[LRZ|VSC|HLRS].sh` which contains:
    - `mpirun -np 192 ./halo_optim.exe < input-optim.txt`
  - Start your batchjob → output file `output_optim.txt`
  - Fill in the table

Note, that the optimization changes the dims-array → modified halo sizes!  
Although halos may be larger, the optimized communication time should be shorter!

Cart_method	Base gridsizes	Block length + stride for each dimension
1	2 2 2	00 00 00
2	2 2 2	00 00 00
1	1 2 4	00 00 00
2	1 2 4	00 00 00
3	1 2 4	00 00 00
0	0 = end	

Halosize/process ~= 26 MB		MPI_Dims_create + MPI_Cart_create		MPIX_Cart_weighted_create(MPIX_WEIGHTS_EQUAL)		MPIX_Cart_weighted_create(...weights...)	
Base grid sizes	Nodes x CPUs x cores	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communication time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]
2 x 2 x 2	_ x _ x _	= baseline	_ = _ x _ x _ _ = _ x _ x _ _ = _ x _ x _	= _____ % of baseline	_ = _ x _ x _ _ = _ x _ x _ _ = _ x _ x _	Reported by MPI_Cart_weighted_create() Same as with MPIX_WEIGHTS_EQUAL	_ = _ x _ x _ _ = _ x _ x _ _ = _ x _ x _



## Exercise: Results – HLRS, Stuttgart, hazelhen

Halosize/process ≈ 26 MB		MPI_Dims_create + MPI_Cart_create		MPIX_Cart_weighted_ create(MPIX_WEIGHTS_EQUAL)		MPIX_Cart_weighted_ create(...weights...)	
Base grid sizes	Nodes x CPUs x cores	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communication time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]
<b>2</b> x <b>2</b> x <b>2</b>	<b>8</b> x 2 x 12	78.748 = baseline	8 = <b>8</b> x 1 x 1 6 = <b>1</b> x 2 x 3 4 = <b>1</b> x 1 x 4	50.971 = 65% of baseline	8 = <b>2</b> x 2 x 2 6 = <b>2</b> x 1 x 3 4 = <b>2</b> x 1 x 2	Same as with MPIX_WEIGHTS_EQUAL	
<b>1</b> x <b>2</b> x <b>4</b>	<b>8</b> x 2 x 12	168.891 = baseline	Same as above	64.691 = 38% of baseline	8 = <b>2</b> x 2 x 2 6 = <b>2</b> x 1 x 3 4 = <b>2</b> x 1 x 2	38.406 = 23% of baseline	4 = <b>1</b> x 2 x 2 6 = <b>2</b> x 1 x 3 8 = <b>4</b> x 1 x 2

## Exercise: Results – LRZ, Garching, ivyMUC

Halosize/process ≈ 26 MB		MPI_Dims_create + MPI_Cart_create		MPIX_Cart_weighted_ create(MPIX_WEIGHTS_EQUAL)		MPIX_Cart_weighted_ create(...weights...)	
Base grid sizes	Nodes x CPUs x cores	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communication time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]
<b>2</b> x <b>2</b> x <b>2</b>	<b>12</b> x 2 x 8	34.814 = baseline	8 = <b>6</b> x 1.3 x 1 6 = <b>1</b> x 1.5 x 2 4 = <b>1</b> x 1 x 4	26.675 = 77% of baseline	6 = <b>3</b> x 1 x 2 8 = <b>2</b> x 2 x 2 4 = <b>2</b> x 1 x 2	Same as with MPIX_WEIGHTS_EQUAL	
<b>1</b> x <b>2</b> x <b>4</b>	<b>12</b> x 2 x 8	54.344 = baseline	Same as above	35.665 = 66% of baseline	6 = <b>3</b> x 1 x 2 8 = <b>2</b> x 2 x 2 4 = <b>2</b> x 1 x 2	22.933 = 42% of baseline	4 = <b>1</b> x 2 x 2 6 = <b>3</b> x 1 x 2 8 = <b>4</b> x 1 x 2

# Exercise: Results – VSC, Vienna, \_\_\_(not yet done)\_\_\_

Halosize/process ≈ 26 MB		MPI_Dims_create + MPI_Cart_create		MPIX_Cart_weighted_ create(MPIX_WEIGHTS_EQUAL)		MPIX_Cart_weighted_ create(...weights...)	
Base grid sizes	Nodes x CPUs x cores	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communication time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]
2 x 2 x 2	12 x 2 x 8	_____ = baseline	8 = 6 x 1.3 x 1 6 = 1 x 1.5 x 2 4 = 1 x 1 x 4	_____ = _____% of baseline	___ = - x - x - ___ = - x - x - ___ = - x - x -	Same as with MPIX_WEIGHTS_EQUAL	
1 x 2 x 4	12 x 2 x 8	_____ = baseline	Same as above	_____ = _____% of baseline	___ = - x - x - ___ = - x - x - ___ = - x - x -	_____ = _____% of baseline	___ = - x - x - ___ = - x - x - ___ = - x - x -

# Exercise: Your result: \_\_\_\_\_

Halosize/process ≈ 26 MB		MPI_Dims_create + MPI_Cart_create		MPIX_Cart_weighted_ create(MPIX_WEIGHTS_EQUAL)		MPIX_Cart_weighted_ create(...weights...)	
Base grid sizes	Nodes x CPUs x cores	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communication time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]	Communicat. time [ms]	dims_ml[d=0] dims_ml[d=1] dims_ml[d=2]
2 x 2 x 2	___ x ___ x ___	_____ = baseline	___ = - x - x - ___ = - x - x - ___ = - x - x -	_____ = _____% of baseline	___ = - x - x - ___ = - x - x - ___ = - x - x -	Same as with MPIX_WEIGHTS_EQUAL	
1 x 2 x 4	___ x ___ x ___	_____ = baseline	Same as above	_____ = _____% of baseline	___ = - x - x - ___ = - x - x - ___ = - x - x -	_____ = _____% of baseline	___ = - x - x - ___ = - x - x - ___ = - x - x -

# References

- [1] Pavan Balaji et al. 2009-2012. Topology awareness in MPI Dims create. <https://github.com/mpi-forum/mpi-forumhistoric/issues/195> Accessed 2018-07-19.

*Another approach using the existing MPI\_Cart\_create() interface:*

- [2] William D. Gropp, Using Node [and Socket] Information to Implement MPI Cartesian Topologies, *Parallel Computing*, 2019, and in: *Proceedings of the 25th European MPI User' Group Meeting, EuroMPI'18*, ACM, New York, NY, USA, 2018, pp. 18:1-18:9. doi:10.1145/3236367.3236377.

Slides: <http://wgropp.cs.illinois.edu/bib/talks/tdata/2018/nodecart-final.pdf>.

*And for unstructured grids:*

- [3] T. Hoefler and M. Snir. 2011. Generic Topology Mapping Strategies for Large-scale Parallel Architectures. In *Proceedings of the 2011 ACM International Conference on Supercomputing (ICS'11)*. ACM, 75–85.

*Used as theoretical basis for implementing MPI\_Dims\_weighted\_create():*

- [4] Jesper Larsson Träff and Felix Donatus Lübbecke. 2015. Specification Guideline Violations by MPI Dims Create. In *Proceedings of the 22nd European MPI Users' Group Meeting (EuroMPI '15)*. ACM, New York, NY, USA, Article 19, 2 pages.

