

ULFM Process Fault Tolerance reading

Aurelien Bouteiller

FT WG

MPI Forum, March 02
Chicago, IL



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE



Info, resources, participate



- Issue Ticket (w/ links to PRs)
- <https://github.com/mmpi-forum/mmpi-issues/issues/20>
- Implementation available
 - Version 1.1 based on Open MPI 1.6 released early November 2015
<https://bitbucket.org/icldistcomp/ulfm>
 - Full communicator-based (point-to-point and all flavors of collectives) support
 - Network support IB, uGNI, TCP, SM
 - Runs with ALPS, PBS, etc...
 - RMA, I/O in progress

<http://fault-tolerance.org/>

Minimal Feature Set for a Resilient MPI

1. Failure Notification
2. Error Propagation
3. Error Recovery

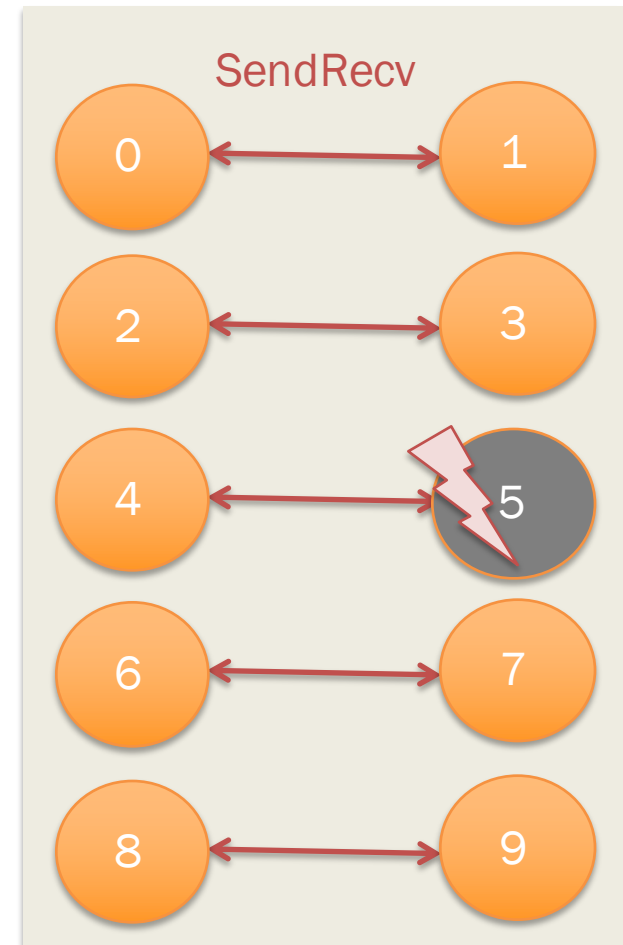
Not all recovery strategies require all of these features, that's why the interface splits notification, propagation and recovery.



ULFM is not a recovery strategy, but a minimalistic set of building blocks for more complex recovery strategies.

Errors are visible only for operations that can't complete

- New error codes to deal with failures
 - **MPI_ERROR_PROC_FAILED**: report that the operation discovered a newly dead process. Returned from all blocking function, and all completion functions.
 - **MPI_ERROR_PROC_FAILED_PENDING**: report that a non-blocking MPI_ANY_SOURCE potential sender has been discovered dead.
 - **MPI_ERROR_REVOKED**: a communicator has been declared improper for further communications. All future communications on this communicator will raise the same error code, with the exception of a handful of recovery functions
- Operations that **can't complete** return **ERR_PROC_FAILED**
 - State of MPI objects unchanged (communicators, etc)
 - Repeating the same operation has the same outcome
- Operations that **can be completed** return **MPI_SUCCESS**
 - Pt-2-pt operations between non failed ranks can continue
- Leverage on existing error handler infrastructure
 - MPI_COMM_SET_ERRHANDLER
 - conveniently capture and manage the new survivable error codes



Example: only rank4 should report the failure of rank 5

Summary of new functions

- `MPI_Comm_failure_ack(comm)`
 - Resumes matching for `MPI_ANY_SOURCE`
- `MPI_Comm_failure_get_acked(comm, &group)`
 - Returns to the user the group of processes acknowledged to have failed

- `MPI_Comm_revoke(comm)`
 - **Non-collective**, interrupts all operations on `comm` (future or active, at all ranks) by raising `MPI_ERR_REVOKED`

- `MPI_Comm_shrink(comm, &newcomm)`
 - Collective, creates a new communicator without failed processes (identical at all ranks)
- `MPI_Comm_agree(comm, &mask)`
 - Collective, agrees on the AND value on binary mask, ignoring failed processes (reliable AllReduce), and the return code

Notification

Propagation

Recovery

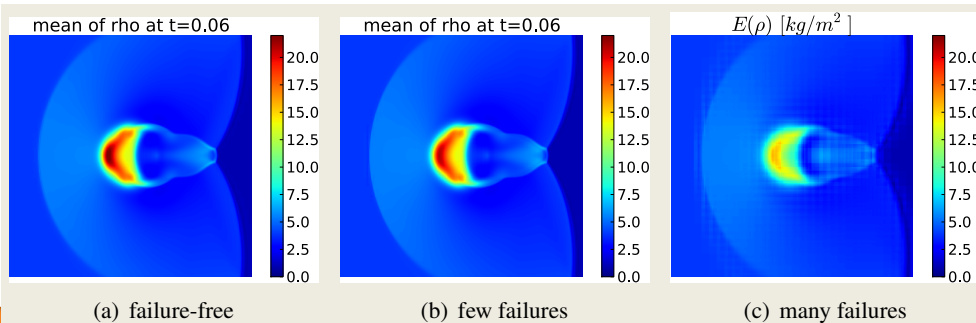
Bibliography of users' activity

These works use ULFM

FRAMEWORKS USING ULFM


LFLR, FENIX, FTLA, Falanx, X10

- HAMOUDA, Sara S., MILTHORPE, Josh, STRAZDINS, Peter E., et al. A Resilient Framework for Iterative Linear Algebra Applications in X10. In : *16th IEEE International Workshop on Parallel and Distributed Scientific and Engineering Computing (PDSEC 2015)*. 2015.
- ST PAULI, P. Arbenz et SCHWAB, Ch. Intrinsic fault tolerance of multi level Monte Carlo methods. *ETH Zurich, Computer Science Department, Tech. Rep.*, 2012.
- PAULI, Stefan, KOHLER, Manuel, et ARBENZ, Peter. A fault tolerant implementation of Multi-Level Monte Carlo methods. In : *PARCO*. 2013. p. 471-480.
- BLAND, Wesley, DU, Peng, BOUTEILLER, Aurelien, et al. Extending the scope of the Checkpoint-on-Failure protocol for forward recovery in standard MPI. *Concurrency and computation: Practice and experience*, 2013, vol. 25, no 17, p. 2381-2393.
- ALI, Md Mortuza, SOUTHERN, James, STRAZDINS, Peter, et al. Application Level Fault Recovery: Using Fault-Tolerant Open MPI in a PDE Solver. In : *Parallel & Distributed Processing Symposium Workshops (IPDPSW), 2014 IEEE International*. IEEE, 2014. p. 1169-1178.
- NAUGHTON, Thomas, ENGELMANN, Christian, VALLÉE, Geoffroy, et al. Supporting the development of resilient message passing applications using simulation. In : *Parallel, Distributed and Network-Based Processing (PDP), 2014 22nd Euromicro International Conference on*. IEEE, 2014. p. 271-278.
- ENGELMANN, Christian et NAUGHTON, Thomas. Improving the Performance of the Extreme-scale Simulator. In : *Proceedings of the 2014 IEEE/ACM 18th International Symposium on Distributed Simulation and Real Time Applications*. IEEE Computer Society, 2014. p. 198-207.
- TERANISHI, Keita et HEROUX, Michael A. Toward Local Failure Local Recovery Resilience Model using MPI-ULFM. In : *Proceedings of the 21st European MPI Users' Group Meeting*. ACM, 2014. p. 51.
- ALI, Md Mohsin, STRAZDINS, Peter E., HARDING, Brendan, et al. A fault-tolerant gyrokinetic plasma application using the sparse grid combination technique. In : *High Performance Computing & Simulation (HPCS), 2015 International Conference on*. IEEE, 2015. p. 499-507.
- VALLÉE, Geoffroy, NAUGHTON, Thomas, BOHM, Swen, et al. A runtime environment for supporting research in resilient HPC system software & tools. In : *Computing and Networking (CANDAR), 2013 First International Symposium on*. IEEE, 2013. p. 213-219.
- ZOUNMEVO, Judicael A., KIMPE, Dries, ROSS, Robert, et al. Extreme-scale computing services over MPI: Experiences, observations and features proposal for next-generation message passing interface. *International Journal of High Performance Computing Applications*, 2014, vol. 28, no 4, p. 435-449.
- NAUGHTON, Thomas, BÖHM, Swen, ENGELMANN, Christian, et al. Using Performance Tools to Support Experiments in HPC Resilience. In : *Euro-Par 2013: Parallel Processing Workshops*. Springer Berlin Heidelberg, 2014. p. 727-736.
- ENGELMANN, Christian et NAUGHTON, Thomas. A NETWORK CONTENTION MODEL FOR THE EXTREME-SCALE SIMULATOR.
- GAMELL, Marc, KATZ, Daniel S., KOLLA, Hemanth, et al. Exploring automatic, online failure recovery for scientific applications at extreme scales. In : *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*. IEEE Press, 2014. p. 895-906.
- XIAOGUANG, Ren, XINHAI, Xu, YUHUA, Tang, et al. An Application-Level Synchronous Checkpoint-Recover Method for Parallel CFD Simulation. In : *Computational Science and Engineering (CSE), 2014 IEEE 17th International Conference on*. IEEE, 2014. p. 58-65.
- Judicael A. Zounmevo, Dries Kimpe, Robert Ross, and Ahmad Afsahi. 2013. Using MPI in high-performance computing services. In *Proceedings of the 20th European MPI Users' Group Meeting (EuroMPI '13)*. ACM, New York, NY, USA, 43-48. SE, 2013 IEEE 16th International Conference on. IEEE, 2013. p. 58-65.
- Junho Ahn, "N Fault-Tolerant Sender-Based Message Logging for Group Communication-Based Message Passing Systems," in *Computational Science and Engineering (CSE), 2014 IEEE 17th International Conference on*, vol., no., pp.1296-1301, 19-21 Dec. 2014.




Credits: ETH Zurich

Figure 5. Results of the FT-MLMC implementation for three different failure scenarios.



The performance improvement due to using ULFM v1.0 for running the LULESH proxy application [3] (a shock hydrodynamics stencil based simulation) running on 64 processes on 16 nodes with

Resolving transitive dependencies



```


proc_failed_err_handler(MPI_Comm comm, int err, ...) {
    if(err == MPI_ERR_PROC_FAILED ||
        err == MPI_ERR_REVOKED) {
        if(err == MPI_ERR_PROC_FAILED) MPI_Comm_revoke(comm);
        recovery(comm);
    }
}

ft_transitive_deps(void) {
    for(i=0; i<nbrecv; i++) {
        if(myrank>0) MPI_Irecv(buff, count, datatype,
                               myrank-1, tag, comm, &req);
        if(myrank<n) MPI_Send(buff2, count, datatype,
                               myrank+1, tag, comm, &req); }
}
    
```

- P1 fails

- P2 raises an error and wants to change comm pattern to do application recovery
- but P3..Pn are stuck in their posted recv
- P2 can unlock them with Revoke
- P3..Pn join P2 in the recovery

Full Capabilities Recovery




- Some applications are moldable
 - Shrink creates a new communicator on which collectives work
- Some applications are not moldable
 - Spawn can recreate a “same size” communicator
 - It is easy to reorder the ranks according to the original ordering

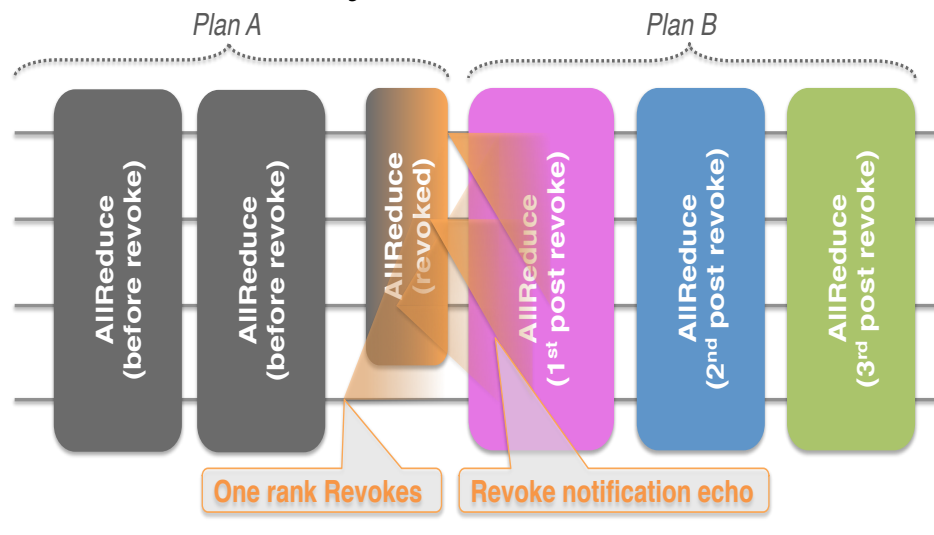
Scalable Resilient Constructs: Revoke

Darter, ugni network, 6000 processes

Revoke Time and Perturbation in Barrier (np=6000)



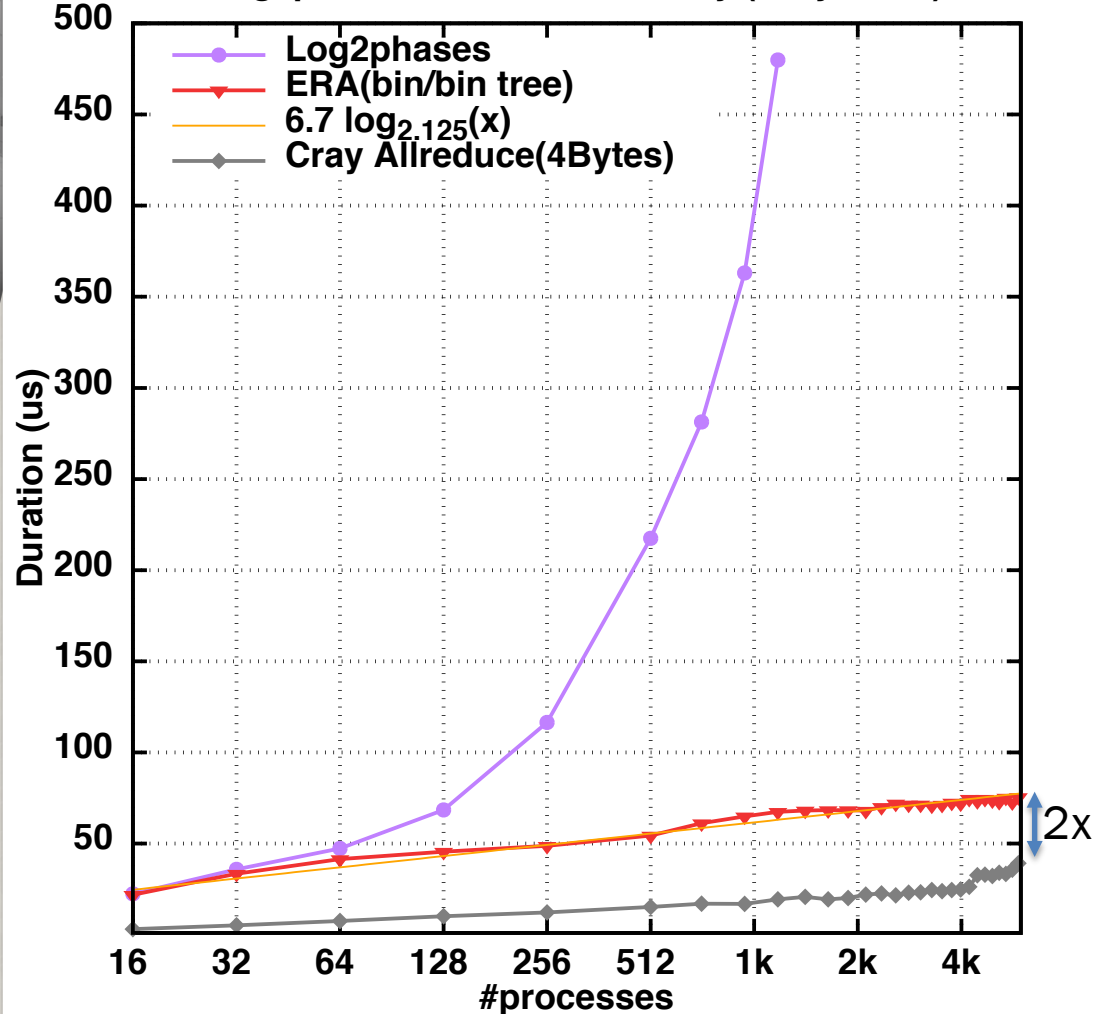
- BMG* Revoke **propagation** in less than **100μs**
- First post-Revoke collective operation sustains some performance degradation resulting from the network jitter associated with the circulation of revoke tokens
- After the fifth Barrier (approximately **700μs**), the Revoke reliable broadcast has **completely terminated**, therefore leaving the application free from observable jitter.



* Bouteiller, A., Bosilca, G., Dongarra, J.J. "Plan B: Interruption of Ongoing MPI Operations to Support Failure Recovery," In *Proceedings of the 22nd European MPI Users' Group Meeting (EuroMPI '15)*. ACM

Scalable Resilient Agreement

Log2phases vs ERA Scalability (Cray XC30)



- Novel Early Returning Agreement algorithm*
- Logarithmic topology & logarithmic computation: scalable
- 2x the Cray AllReduce latency at 6k processors!

* Hernaut, T., Bouteiller, A., Bosilca, G., Gamell, M., Teranishi, K., Parashar, M., Dongarra, J. "Practical Scalable Consensus for Pseudo-Synchronous Distributed Systems," SuperComputing, Austin, TX, November, 2015

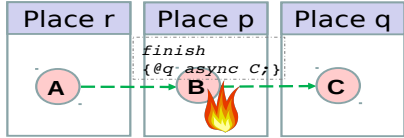
User projects: Resilient X10

- X10 is a PGAS programming language
 - Legacy resilient X10 TCP based

Happens Before Invariance Principle (HBI):

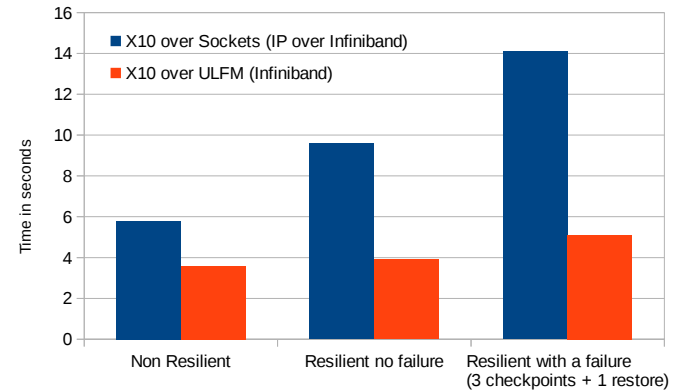
Failure of a place should not alter the happens before relationship between statements at the remaining places.

```
try{ /*Task A*/  
  at (p) { /*Task B*/  
    finish { at (q) async { /*Task C*/ } }  
  }  
} catch(dpe:DeadPlaceException){ /*recovery steps*/}  
D;
```



By applying the HBI principle, Resilient X10 will ensure that statement D executes after Task C finishes, despite the loss of the synchronization construct (finish) at place p

- MPI operations in resilient X10 runtime
 - Progress loop does MPI_Iprobe, post needed recv according to probes
 - Asynchronous background collective operations (on multiple different comms to form 2d grids, etc).
- Recovery
 - Upon failure, all communicators recreated (from shrinking a large communicator with spares, or using MPI_COMM_SPAWN to get new ones)
 - Ranks reassigned identically to rebuild the same X10 “teams”
- Injection of FT layer
 - Unnecessary, x10 has a runtime that hides all MPI from the application and handles failures internally



The performance improvement due to using ULMF v1.0 for running the LULESH proxy application [3] (a shock hydrodynamics stencil based simulation) running on 64 processes on 16 nodes with

User projects: Fenix+S3D

- Fenix is a framework to provide scoped user level checkpoint/restart
 - Provides some of the same services provided by the “MPI_Reinit” idea floated around by T. Gamblin
 - Recover failed processes with revoke-shrink-spawn-reorder sequence
 - Recovered and surviving processes jump back to the start (longjump in Fenix_init)
 - Fenix has helpers to perform user directed “in-memory” or “buddy” checkpointing (and reload)
 - Injection of FT layer: PMPI based
- **Fenix_Checkpoint_Allocate** mark a memory segment (baseptr,size) as part of the checkpoint.
- **Fenix_Init** Initialize Fenix, and restart point after a recovery, status contains info about the restart mode
- **Fenix_Comm_Add** can be used to notify Fenix about the creation of user communicators
- **Fenix_Checkpoint** performs a checkpoint of marked segments

```
1 allocate(yzpc(nx,ny,nz,nslvs))
2 allocate(other_arrays)
3 call MPI_Init()
4 [...] ! Initialize non-conflicting modules
5 call Fenix_Checkpoint_Allocate(C_LOC(yzpc),
6     sizeof(yzpc),ckpt_yzpc)
7 call Fenix_Init(Fenix_Neighbors,PEER_NODE_SIZE,
8     Fenix_resume_to_init, status, C_LOC(world))
9
10 if(status.eq.Fenix_st_survivor) then
11     [...] ! Finalize conflicting modules
12 endif
13 [...] ! Initialize conflicting modules
14 if(status.eq.Fenix_st_new)
15     call initialize_yzpc()
16 endif
17
18 do ! Main loop
19     [...] ! Iterate and update yzpc array
20     if(mod(step-1,CHECKPOINT_PERIOD).eq.0) then
21         call Fenix_Checkpoint(ckpt_yzpc);
22     endif
23 enddo
24
25 call Fenix_Finalize()
26 call MPI_Finalize()
```

User projects: Fenix+S3D

- S3D is a production, highly parallel method-of-lines solver for PDEs
 - used to perform first-principles-based direct numerical simulations of turbulent combustion
- S3D rendered fault tolerant using Fenix/ULFM
- 35 lines of code modified in S3D in total!
- Order of magnitude performance improvement in failure scenarios
 - thanks to online recovery and in-memory checkpoint advantage over I/O based checkpointing
- Injection of FT layer: addition of a couple of Fenix calls

```
1 call MPI_Comm_split(gcomm, py+1000*pz, r, xcomm)
2 call MPI_Comm_split(gcomm, px+1000*pz, r, ycomm)
3 call MPI_Comm_split(gcomm, px+1000*py, r, zcomm)
4 call Fenix_Comm_Add(xcomm);
5 call Fenix_Comm_Add(ycomm);
6 call Fenix_Comm_Add(zcomm);
7 [...]
8 call MPI_Comm_split(gcomm, xid, r, yz_comm)
9 call MPI_Comm_split(gcomm, yid, r, xz_comm)
10 call MPI_Comm_split(gcomm, zid, r, xy_comm)
11 call Fenix_Comm_Add(yz_comm);
12 call Fenix_Comm_Add(xz_comm);
13 call Fenix_Comm_Add(xy_comm);
```

S3D Code snippet to declare to Fenix the communicators to recover

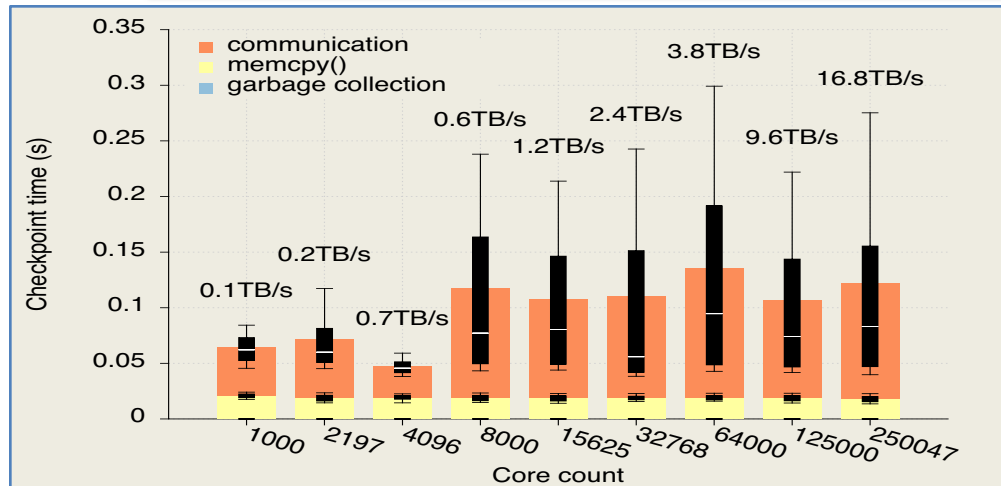


Fig. 3. Checkpoint time for different core counts (8.6 MB/core). The numbers above each test show the aggregated bandwidth (the total checkpoint size over the average checkpoint time).